

# EduSphere API Contracts — GraphQL Federation Schema

**Reference document for Claude Code and all developers.** Every subgraph, entity, query, mutation, subscription, input type, and directive is defined here. Do NOT deviate from these definitions without updating this document first. This document is the **single source of truth** for the EduSphere supergraph API surface.

---

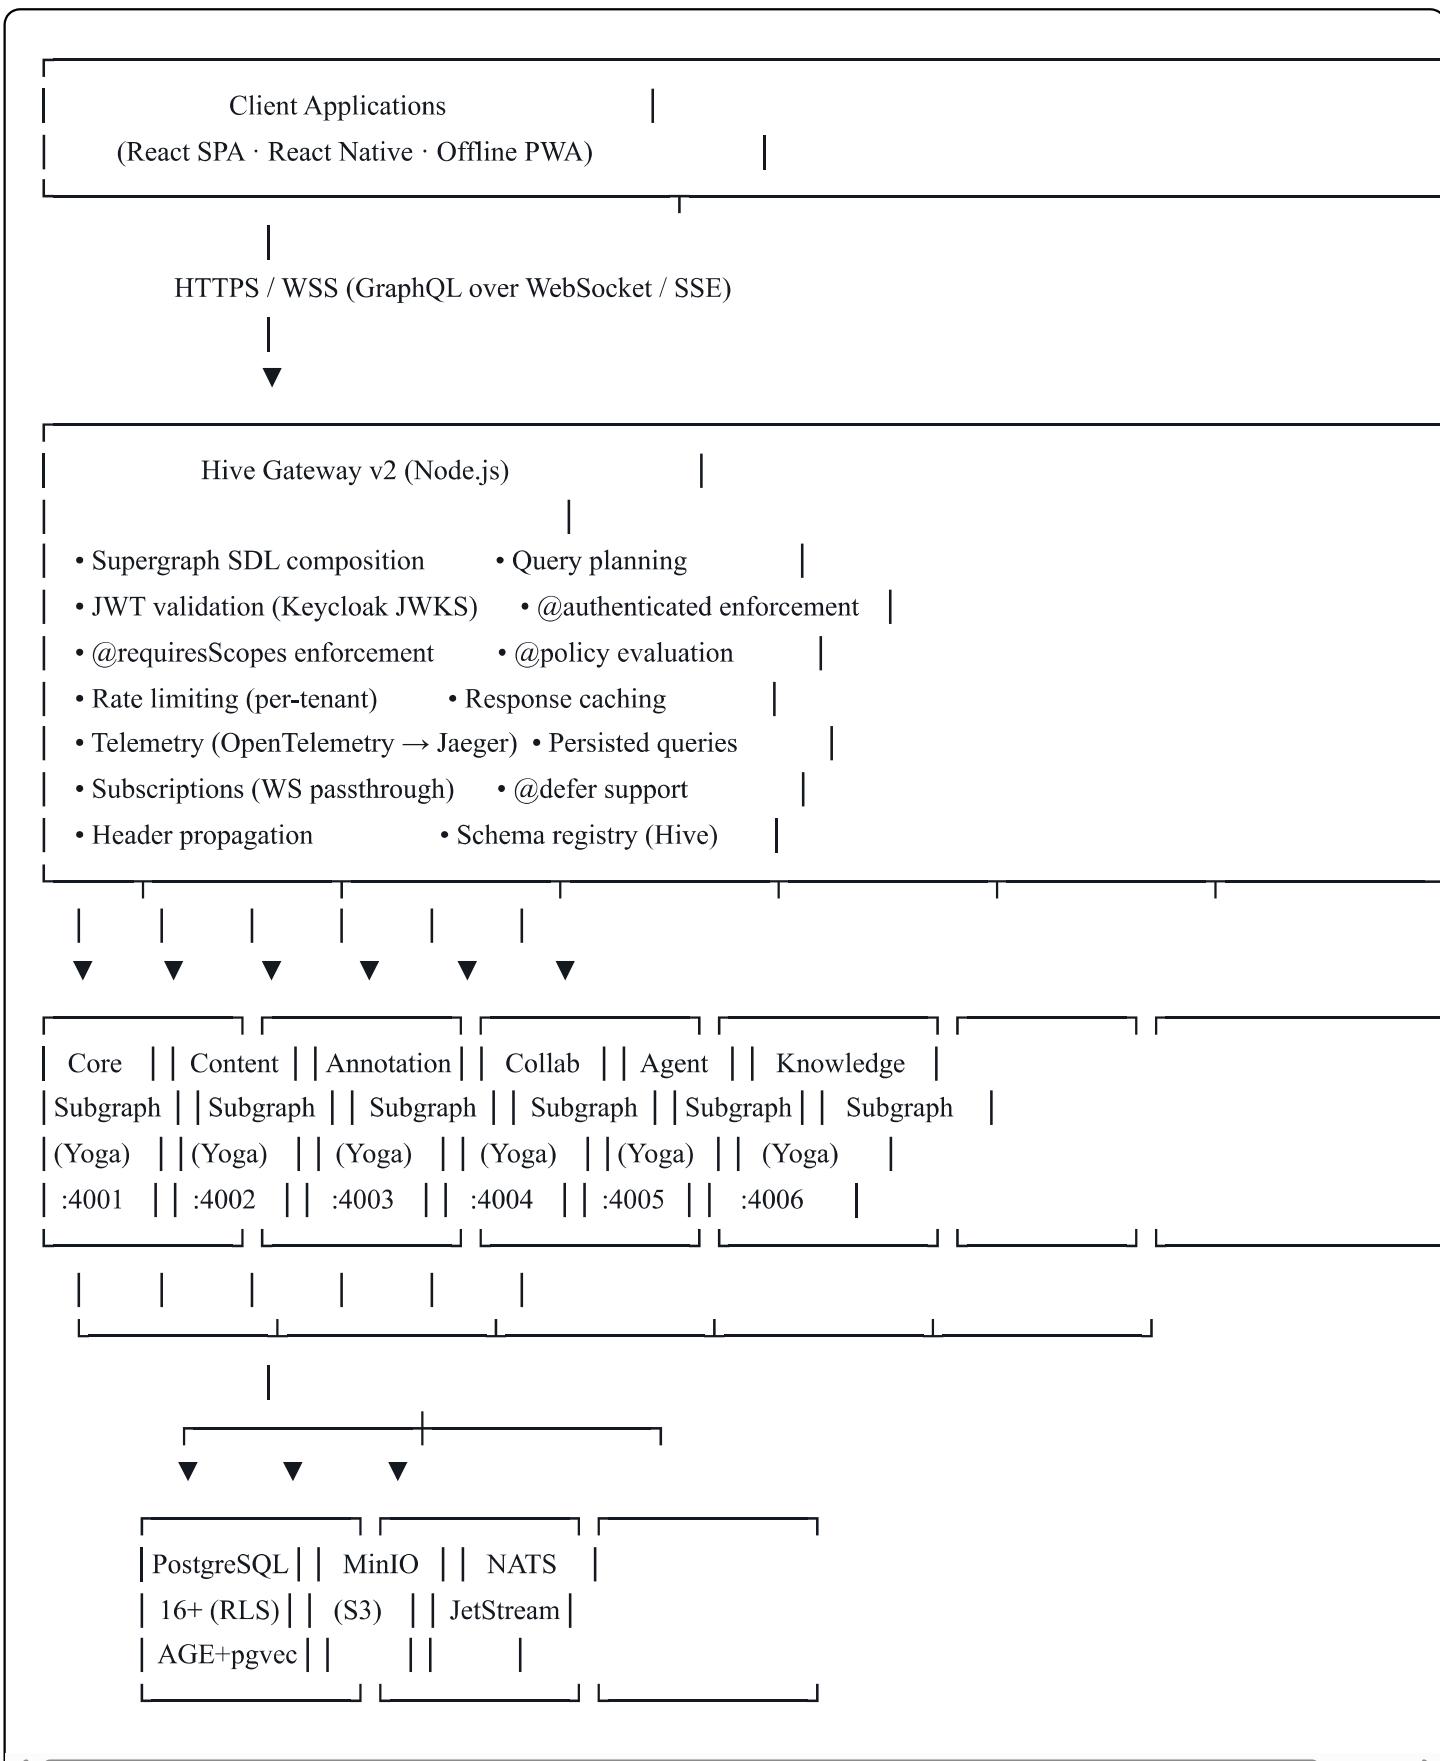
## Table of Contents

1. [Federation Architecture Overview](#)
  2. [Subgraph Decomposition Map](#)
  3. [Shared Types & Scalars](#)
  4. [Pagination & Connection Spec](#)
  5. [Authentication & Authorization Directives](#)
  6. [Error Handling Contract](#)
  7. [Subgraph: Core \(Tenants & Users\)](#)
  8. [Subgraph: Content \(Courses, Media, Transcripts\)](#)
  9. [Subgraph: Annotation \(Layers, Sketches, Comments\)](#)
  10. [Subgraph: Collaboration \(CRDT, Real-time\)](#)
  11. [Subgraph: Agent \(Definitions, Executions, Templates\)](#)
  12. [Subgraph: Knowledge \(Graph, Embeddings, Semantic Search\)](#)
  13. [Subscriptions Contract](#)
  14. [File Upload Contract \(Multipart\)](#)
  15. [Gateway Configuration & Router Settings](#)
  16. [Schema Governance & Evolution Rules](#)
  17. [Client Operation Examples](#)
  18. [TypeScript Codegen Configuration](#)
  19. [AI/ML Architecture Integration](#)
  20. [Technology Stack Cross-Reference](#)
- 

## 1. Federation Architecture Overview

EduSphere uses **GraphQL Federation v2.7+** (Apollo Federation v2 spec compatible) with **Hive Gateway v2** (MIT-licensed) as the supergraph gateway. Each subgraph is a standalone NestJS service backed by **GraphQL Yoga** with `(YogaFederationDriver)` from `@graphql-yoga/nestjs`.

**Why Hive Gateway over Apollo Router?** Apollo Router uses the Elastic License v2 (ELv2), which restricts offering it as a managed service. Hive Gateway v2 is fully MIT-licensed, supports the same Federation v2.7 spec, and is nearly 2x faster than competitors among JS-based gateways. It runs on all runtimes (Node, Deno, Bun, Cloudflare Workers). See EduSphere Architecture Guide §5 for details.



## Key Decisions:

- **Schema-first design:** SDL files are the source of truth; resolvers implement the contract

- **MIT-licensed gateway:** Hive Gateway v2 (MIT) over Apollo Router (ELv2) — full Federation v2 spec compatibility
  - **GraphQL Yoga subgraphs:** Each NestJS service uses `@graphql-yoga/nestjs` with `YogaFederationDriver` for federation
  - **Transport-based multi-tenancy:** JWT carries `tenant_id`; gateway propagates `x-tenant-id` header to subgraphs; RLS enforces at database level
  - **Relay Cursor Connections:** all list fields use the Relay connection spec for forward/backward cursor pagination
  - **Demand-oriented schema:** types are shaped for client needs, not database tables
  - **Entity ownership:** each entity is owned by exactly one subgraph; other subgraphs extend with `@key` stubs
  - **Soft deletes hidden:** `deleted_at` is never exposed; filtered out in resolvers
  - **Three-layer AI architecture:** Vercel AI SDK (LLM abstraction) → LangGraph.js (agent orchestration) → LlamaIndex.TS (RAG/knowledge)
  - **HybridRAG pattern:** pgvector semantic search + Apache AGE graph traversal fused before LLM generation
  - **Hive Schema Registry:** Breaking change detection and schema evolution tracking via GraphQL Hive
- 

## 2. Subgraph Decomposition Map

Subgraph	Port	Owned Entities	Extended Entities	Domain
Core	4001	<code>Tenant</code> , <code>User</code>	—	Identity, tenancy, auth
Content	4002	<code>Course</code> , <code>Module</code> , <code>MediaAsset</code> , <code>Transcript</code> , <code>TranscriptSegment</code>	<code>User</code> ( <code>createdCourses</code> ), <code>Tenant</code>	Courses, media pipeline, transcription
Annotation	4003	<code>Annotation</code>	<code>User</code> , <code>MediaAsset</code>	Markings, sketches, spatial comments
Collaboration	4004	<code>CollabDocument</code> , <code>CollabSession</code>	<code>User</code> , <code>Annotation</code>	CRDT persistence, real-time presence
Agent	4005	<code>AgentDefinition</code> , <code>AgentExecution</code>	<code>User</code> , <code>Annotation</code> , <code>Course</code>	AI agents, templates, executions

Subgraph	Port	Owned Entities	Extended Entities	Domain
Knowledge	4006	Concept, Person, Term, Source, TopicCluster, KnowledgeRelation, SemanticSearchResult	MediaAsset, TranscriptSegment, Annotation	Knowledge graph, embeddings, search

## Entity Ownership Rules

Entity	Owning Subgraph	@key fields
Tenant	Core	id
User	Core	id
Course	Content	id
Module	Content	id
MediaAsset	Content	id
Transcript	Content	id
TranscriptSegment	Content	id
Annotation	Annotation	id
CollabDocument	Collaboration	id
CollabSession	Collaboration	id
AgentDefinition	Agent	id
AgentExecution	Agent	id
Concept	Knowledge	id
Person	Knowledge	id
Term	Knowledge	id
Source	Knowledge	id
TopicCluster	Knowledge	id

## 3. Shared Types & Scalars

All subgraphs MUST import the following shared definitions. These are declared in a shared package `@edusphere/graphql-shared` and included via build-time codegen.

graphql

```
# ----- File: packages/graphql-shared/src/scalars.graphql -----
```

```
"""ISO 8601 date-time string with timezone (e.g. "2025-01-15T10:30:00Z")"""
scalar DateTime
```

```
"""UUID v4 string (e.g. "550e8400-e29b-41d4-a716-446655440000")"""
scalar UUID
```

```
"""Arbitrary JSON object"""
scalar JSON
```

```
"""Arbitrary JSON object (alias for interoperability)"""
scalar JSONObject
```

```
"""Base64-encoded binary data"""
scalar Base64
```

```
"""

```

```
Opaque cursor string for Relay-style pagination.
```

```
Clients MUST NOT parse or construct cursors — treat as opaque tokens.
```

```
"""

```

```
scalar Cursor
```

```
"""URL string (validated as a valid URL)"""
scalar URL
```

```
"""Email address string (validated)"""
scalar EmailAddress
```

```
"""Non-negative integer ( $\geq 0$ )"""
scalar NonNegativeInt
```

```
"""Positive integer ( $> 0$ )"""
scalar PositiveInt
```

```
"""Float between 0.0 and 1.0 inclusive"""
scalar UnitFloat
```

## Shared Enums

```
graphql
```

```
# ----- File: packages/graphql-shared/src/enums.graphql -----
```

```
enum UserRole {  
    SUPER_ADMIN  
    ORG_ADMIN  
    INSTRUCTOR  
    STUDENT  
    RESEARCHER  
}
```

```
enum MediaType {  
    VIDEO  
    AUDIO  
    PDF  
    IMAGE  
    DOCUMENT  
}
```

```
enum TranscriptionStatus {  
    PENDING  
    PROCESSING  
    COMPLETED  
    FAILED  
}
```

```
enum AnnotationType {  
    TEXT  
    SKETCH  
    LINK  
    BOOKMARK  
    SPATIAL_COMMENT  
}
```

```
enum AnnotationLayer {  
    PERSONAL  
    SHARED  
    INSTRUCTOR  
    AI_GENERATED  
}
```

```
enum AgentTemplate {  
    CHAVRUTA  
    SUMMARIZER  
    QUIZ_MASTER  
    RESEARCH_SCOUT  
    EXPLAINER  
}
```

```

CUSTOM
}

enum AgentTrigger {
MANUAL
ON_ANNOTATION
ON_NEW_CONTENT
SCHEDULED
}

enum AgentOutputFormat {
CHAT
OVERLAY
SIDE_PANEL
NOTIFICATION
}

enum ExecutionStatus {
QUEUED
RUNNING
COMPLETED
FAILED
CANCELLED
}

enum TenantPlan {
FREE
STARTER
PROFESSIONAL
ENTERPRISE
}

"""Sort direction for ordered queries"""
enum SortDirection {
ASC
DESC
}

```

## 4. Pagination & Connection Spec

All list fields MUST use the **Relay Cursor Connections Specification**. No raw lists for user-facing collections.

### Connection Types (Generic Pattern)

```
# —— File: packages/graphql-shared/src/pagination.graphql ——
```

```
"""
```

Relay-compliant page info for cursor-based pagination.

Used by all Connection types across every subgraph.

```
"""\n\n
```

```
type PageInfo {
```

```
    """When paginating forwards, are there more items?"""
    hasNextPage: Boolean!
```

```
    """When paginating backwards, are there more items?"""
    hasPreviousPage: Boolean!
```

```
    """Cursor of the first edge in the current page"""
    startCursor: Cursor
```

```
    """Cursor of the last edge in the current page"""
    endCursor: Cursor
```

```
}
```

```
"""\n\n
```

Standard pagination arguments accepted by all connection fields.

At least one of (first, last) MUST be provided.

first + after = forward pagination.

last + before = backward pagination.

Providing both first AND last is STRONGLY DISCOURAGED.

```
"""\n\n
```

```
input ConnectionArgs {
```

```
    """Returns the first N edges after the cursor"""
    first: PositiveInt
```

```
    """Returns edges after this cursor"""
    after: Cursor
```

```
    """Returns the last N edges before the cursor"""
    last: PositiveInt
```

```
    """Returns edges before this cursor"""
    before: Cursor
```

```
}
```

## Naming Convention for Connections

Every entity `(Foo)` generates:

- `(FooConnection)` — contains `(edges)` and `(pageInfo)`
- `(FooEdge)` — contains `(node)` and `(cursor)`

graphql

```

# Example: CourseConnection
type CourseConnection {
  edges: [CourseEdge!]!
  pageInfo: PageInfo!
  """Total count of items matching filters (computed only if requested)"""
  totalCount: NonNegativeInt
}

type CourseEdge {
  node: Course!
  cursor: Cursor!
}

```

## Pagination Limits

Context	Default <code>first</code>	Max <code>first</code> / <code>last</code>
Standard lists	20	100
Transcript segments	50	500
Search results	10	50
Agent executions	20	100
Knowledge graph relations	20	200

If `first`/`last` exceeds the maximum, the server clamps to max and returns a `PAGINATION_CLAMPED` extension warning.

## 5. Authentication & Authorization Directives

### Federation v2 Auth Directives

```
graphql
```

# —— Applied at the gateway level (Hive Gateway config) ——

----

Requires a valid JWT. Unauthenticated requests receive 401.

Applied to all fields by default; opt-out with @public.

----

directive `@authenticated` on FIELD\_DEFINITION | OBJECT | INTERFACE | SCALAR | ENUM

----

Requires specific OAuth2 scopes from the JWT.

Scopes use AND within inner array, OR across outer array.

----

directive `@requiresScopes`(

  scopes: [[String!]!]!

) on FIELD\_DEFINITION | OBJECT | INTERFACE | SCALAR | ENUM

----

Custom authorization policy evaluated by the router.

Policies are defined in router.yaml and evaluated against JWT claims.

----

directive `@policy`(

  policies: [[String!]!]!

) on FIELD\_DEFINITION | OBJECT | INTERFACE | SCALAR | ENUM

## Custom Application Directives

graphql

```
# —— File: packages/graphql-shared/src/directives.graphql ——
```

```
"""
```

Marks a field or type as publicly accessible without authentication.

Overrides the default @authenticated requirement.

```
"""
```

```
directive @public on FIELD_DEFINITION | OBJECT
```

```
"""
```

Requires the caller to have one of the specified roles.

Evaluated in subgraph resolvers using context.user.role.

```
"""
```

```
directive @requiresRole(
```

```
  roles: [UserRole!]!
```

```
) on FIELD_DEFINITION | OBJECT
```

```
"""
```

Marks a field as only accessible by the resource owner.

The resolver checks context.user.id against the entity's userId/creatorId.

```
"""
```

```
directive @ownerOnly on FIELD_DEFINITION
```

```
"""
```

Rate-limits a field to N calls per window per user.

Evaluated at the router level via the rate-limiter plugin.

```
"""
```

```
directive @rateLimit(
```

```
  max: Int!
```

```
  window: String! # e.g., "1m", "1h", "1d"
```

```
) on FIELD_DEFINITION
```

```
"""
```

Marks a field as deprecated with a reason and optional removal date.

Clients using deprecated fields will see warnings in tooling.

```
"""
```

```
directive @deprecated(
```

```
  reason: String!
```

```
  removalDate: String # ISO 8601 date
```

```
) on FIELD_DEFINITION | ENUM_VALUE | ARGUMENT_DEFINITION | INPUT_FIELD_DEFINITION
```

## Role-to-Scope Mapping

Role	JWT Scopes	Capabilities
SUPER_ADMIN	admin:*	Full system access, cross-tenant

Role	JWT Scopes	Capabilities
ORG_ADMIN	[org:manage], [org:users], [org:billing]	Manage tenant settings, users, billing
INSTRUCTOR	[course:write], [annotation:write], [agent:write], [media:upload]	Create/edit courses, annotations, agents
STUDENT	[course:read], [annotation:write:own], [agent:execute]	View courses, create own annotations, run agents
RESEARCHER	[course:read], [annotation:read], [knowledge:read], [knowledge:write]	Read content, query knowledge graph, add concepts

## JWT Claims Structure (Propagated by Router)

typescript

```
// Extracted from Keycloak JWT and propagated as headers
interface JWTClaims {
  sub: string; // Keycloak user ID
  tenant_id: string; // UUID of the tenant
  user_id: string; // EduSphere internal user UUID
  email: string;
  roles: UserRole[];
  scopes: string[];
  iat: number;
  exp: number;
}
```

## Header Propagation (Router → Subgraphs)

yaml

```
# Headers the router propagates to all subgraph requests
headers:
  all:
    request:
      - propagate:
          named: "x-tenant-id"
      - propagate:
          named: "x-user-id"
      - propagate:
          named: "x-user-role"
      - propagate:
          named: "x-user-scopes"
      - propagate:
          named: "authorization"
      - propagate:
          named: "x-request-id"
      - propagate:
          named: "x-correlation-id"
```

## 6. Error Handling Contract

### Standard Error Format

All subgraphs MUST return errors following this structure. The `extensions` field provides machine-readable metadata for client error handling.

```
graphql
```

```
# Conceptual — errors are part of the GraphQL response, not the schema
# Included here as the contract all resolvers MUST follow
```

```
"""
```

```
Standard error extension fields returned in every GraphQL error.
```

```
"""
```

```
type ErrorExtension {
    """Machine-readable error code (UPPER_SNAKE_CASE)"""
    code: String!
    """HTTP-equivalent status code"""
    status: Int!
    """ISO 8601 timestamp of the error"""
    timestamp: DateTime!
    """Correlation ID for tracing"""
    requestId: String!
    """Field path that caused the error"""
    path: [String!]
    """Additional context for debugging (only in development)"""
    debug: JSON
}
```

## Error Codes Catalog

Code	HTTP Status	Description	Retryable
UNAUTHENTICATED	401	Missing or invalid JWT	No
FORBIDDEN	403	Valid JWT but insufficient permissions	No
NOT_FOUND	404	Requested resource does not exist	No
CONFLICT	409	Resource version conflict (optimistic locking)	Yes
VALIDATION_ERROR	400	Input validation failure	No
RATE_LIMITED	429	Too many requests	Yes (after <a href="#">retryAfter</a> )
TENANT_QUOTA_EXCEEDED	403	Tenant has exceeded plan limits	No
MEDIA_PROCESSING_ERROR	500	Media transcoding/transcription failure	Yes
AGENT_EXECUTION_ERROR	500	Agent runtime failure	Yes
GRAPH_QUERY_ERROR	500	Apache AGE query failure	Yes

Code	HTTP Status	Description	Retryable
EMBEDDING_ERROR	500	Vector embedding generation failure	Yes
INTERNAL_ERROR	500	Unhandled server error	Yes
PAGINATION_CLAMPED	200	Requested page size was clamped (warning)	N/A
OPTIMISTIC_LOCK_FAILED	409	Version mismatch on update	Yes

## Error Response Example

```
json
{
  "data": null,
  "errors": [
    {
      "message": "Course not found",
      "locations": [{ "line": 2, "column": 3 }],
      "path": ["course"],
      "extensions": {
        "code": "NOT_FOUND",
        "status": 404,
        "timestamp": "2025-06-15T10:30:00Z",
        "requestId": "req_abc123def456"
      }
    }
  ]
}
```

## Partial Success Pattern

For mutations that operate on multiple items, subgraphs MUST return a `MutationResult` union indicating per-item success/failure:

```
graphql
```

```
type MutationSuccess {
  message: String
}

type MutationError {
  message: String!
  code: String!
  field: String
}

union MutationResult = MutationSuccess | MutationError
```

## 7. Subgraph: Core (Tenants & Users)

**Owns:** Tenant, User **Port:** 4001 **Database tables:** tenants, users

### Schema Definition

**Note:** The `@link(url: "https://specs.apollo.dev/federation/v2.7" ...)` directive imports the Federation v2 specification, which is an open standard. Hive Gateway v2 fully supports this spec — the `specs.apollo.dev` URL is the canonical spec location, not a dependency on Apollo products.

graphql

```
# ----- File: apps/subgraph-core/src/schema.graphql -----
```

```
extend schema
```

```
@link(url: "https://specs.apollo.dev/federation/v2.7",
```

```
import: [
```

```
  "@key",
```

```
  "@shareable",
```

```
  "@inaccessible",
```

```
  "@authenticated",
```

```
  "@requiresScopes",
```

```
  "@policy",
```

```
  "@tag",
```

```
  "@override"
```

```
])
```

```
#-----
```

```
# Tenant
```

```
#-----
```

```
""""
```

```
Top-level organizational unit. All data is tenant-isolated.
```

```
Each tenant represents an educational institution, organization, or team.
```

```
""""
```

```
type Tenant @key(fields: "id") {
```

```
  id: UUID!
```

```
    """Human-readable name of the organization"""
```

```
  name: String!
```

```
    """URL-friendly unique slug (e.g., "bar-ilan-university")"""
```

```
  slug: String!
```

```
    """Current subscription plan"""
```

```
  plan: TenantPlan!
```

```
    """Whether the tenant is active (deactivated tenants cannot access the system)"""
```

```
  isActive: Boolean!
```

```
    """Maximum number of users allowed under current plan"""
```

```
  maxUsers: NonNegativeInt!
```

```
    """Tenant-specific settings and branding"""
```

```
  settings: TenantSettings
```

```
  createdAt: DateTime!
```

```
  updatedAt: DateTime!
```

```
# ----- Resolved by other subgraphs -----
```

```
# users: UserConnection! → resolved by Core itself
```

```
# courses: CourseConnection! → resolved by Content subgraph
```

```
}
```

```
type TenantSettings {
```

```
defaultLanguage: String  
allowedLlmProviders: [String!]  
maxStorageGb: NonNegativeInt  
customOntologyFields: [OntologyFieldDef!]  
brandingColors: BrandingColors  
}
```

```
type OntologyFieldDef {  
  name: String!  
  type: String!  
}
```

```
type BrandingColors {  
  primary: String!  
  secondary: String!  
}
```

```
# -----  
# User  
# -----
```

\*\*\*\*

An authenticated user within a tenant.

Users are synced from Keycloak on first login.

\*\*\*\*

```
type User @key(fields: "id") {  
  id: UUID!  
  """Email address (unique within tenant)"""  
  email: EmailAddress!  
  """Display name shown in UI and collaboration features"""  
  displayName: String!  
  """URL to avatar image"""  
  avatarUrl: URL  
  """Role within the tenant"""  
  role: UserRole!  
  """Whether the user can currently access the system"""  
  isActive: Boolean!  
  """User preferences for UI and behavior"""  
  preferences: UserPreferences  
  createdAt: DateTime!  
  updatedAt: DateTime!
```

```
# ----- Tenant relationship -----  
tenant: Tenant!
```

```
# ----- Resolved by other subgraphs -----  
# annotations: AnnotationConnection! → Annotation subgraph
```

```
# agentDefinitions: AgentDefinitionConnection! → Agent subgraph
# agentExecutions: AgentExecutionConnection! → Agent subgraph
# createdCourses: CourseConnection! → Content subgraph
}
```

```
type UserPreferences {
  language: String
  theme: ThemePreference
  defaultAnnotationLayer: AnnotationLayer
  notificationsEnabled: Boolean
  agentDefaults: AgentPreferenceDefaults
}
```

```
enum ThemePreference {
  LIGHT
  DARK
  SYSTEM
}
```

```
type AgentPreferenceDefaults {
  difficulty: String
  personality: String
}
```

```
# -----
# Queries
# -----
```

```
type Query {
  """
  Returns the currently authenticated user.
  This is the primary entry point for client initialization.
  """
  me: User! @authenticated

  """
  Fetch a user by ID within the current tenant
  """
  user(id: UUID!): User @authenticated
```

```
"""
List all users within the current tenant.
```

```
Filterable by role and activity status.
"""

users(
```

```
  first: PositiveInt = 20
```

```
  after: Cursor
```

```
  last: PositiveInt
```

```
  before: Cursor
```

```
filter: UserFilterInput
orderBy: UserOrderByInput
): UserConnection! @authenticated

"""
Fetch the current tenant's details.

currentTenant: Tenant! @authenticated

"""
Fetch a tenant by slug. Used for SSO login flow to resolve tenant.

tenantBySlug(slug: String!): Tenant @public
}
```

```
# -----  
# Mutations  
# -----
```

```
type Mutation {  
"""
Update the current user's profile and preferences.  
Users can only update their own profile.  
"  
updateMyProfile(input: UpdateProfileInput!): User! @authenticated  
"  
"""
Update a user's role. Requires ORG_ADMIN or SUPER_ADMIN.  
"  
updateUserRole(  
  userId: UUID!  
  role: UserRole!  
): User! @authenticated @requiresScopes(scopes: ["org:users"])  
"  
"
```

```
"""
Deactivate a user (soft). Requires ORG_ADMIN.  
Deactivated users cannot log in.  
"  
deactivateUser(  
  userId: UUID!  
): User! @authenticated @requiresScopes(scopes: ["org:users"])  
"  
"
```

```
"""
Reactivate a previously deactivated user.  
"  
reactivateUser(  
  userId: UUID!
```

```
): User! @authenticated @requiresScopes(scopes: ["org:users"])
```

```
""""
```

```
Update tenant settings. Requires ORG_ADMIN.
```

```
""""
```

```
updateTenantSettings(
```

```
    input: UpdateTenantSettingsInput!
```

```
): Tenant! @authenticated @requiresScopes(scopes: ["org:manage"])
```

```
}
```

---

```
# -----  
# Input Types  
# -----
```

```
input UpdateProfileInput {
```

```
    displayName: String
```

```
    avatarUrl: URL
```

```
    preferences: UserPreferencesInput
```

```
}
```

```
input UserPreferencesInput {
```

```
    language: String
```

```
    theme: ThemePreference
```

```
    defaultAnnotationLayer: AnnotationLayer
```

```
    notificationsEnabled: Boolean
```

```
    agentDefaults: AgentPreferenceDefaultsInput
```

```
}
```

```
input AgentPreferenceDefaultsInput {
```

```
    difficulty: String
```

```
    personality: String
```

```
}
```

```
input UpdateTenantSettingsInput {
```

```
    defaultLanguage: String
```

```
    allowedLlmProviders: [String!]
```

```
    maxStorageGb: NonNegativeInt
```

```
    brandingColors: BrandingColorsInput
```

```
    customOntologyFields: [OntologyFieldDefInput!]
```

```
}
```

```
input BrandingColorsInput {
```

```
    primary: String!
```

```
    secondary: String!
```

```
}
```

```
input OntologyFieldDefInput {
```

```

name: String!
type: String!
}

input UserFilterInput {
    """Filter by role"""
role: UserRole
"""Filter by active status"""
isActive: Boolean
"""Full-text search on displayName and email"""
search: String
}

input UserOrderByInput {
    field: UserSortField!
    direction: SortDirection!
}

enum UserSortField {
    DISPLAY_NAME
    EMAIL
    CREATED_AT
    ROLE
}

# -----
# Connection Types
# -----

```

```

type UserConnection {
    edges: [UserEdge!]!
    pageInfo: PageInfo!
    totalCount: NonNegativeInt
}

type UserEdge {
    node: User!
    cursor: Cursor!
}

```

## 8. Subgraph: Content (Courses, Media, Transcripts)

**Owns:** Course, Module, MediaAsset, Transcript, TranscriptSegment **Port:** 4002 **Database tables:** courses, modules, media\_assets, transcripts, transcript\_segments

## Schema Definition

graphql

```
# ----- File: apps/subgraph-content/src/schema.graphql -----
```

```
extend schema
```

```
@link(url: "https://specs.apollo.dev/federation/v2.7",
  import: [
    "@key", "@shareable", "@external", "@provides",
    "@requires", "@authenticated", "@requiresScopes", "@tag"
  ]
)
```

---

```
# Entity Stubs (owned by other subgraphs)
```

---

```
type User @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
type Tenant @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

---

```
# Extend User with content fields
#-----
```

```
extend type User {
  """Courses created by this user (instructors/admins only)"""
  createdCourses(
    first: PositiveInt = 20
    after: Cursor
    filter: CourseFilterInput
  ): CourseConnection!
}
```

---

```
#-----
```

---

```
# Course
#-----
```

```
"""
```

A course is the primary organizational unit for content.

Courses contain modules, which contain media assets.

Courses support versioning and can be forked.

```
"""
```

```
type Course @key(fields: "id") {
  id: UUID!
  title: String!
```

```
description: String
"""Structured syllabus as a JSON tree"""
syllabusJson: JSON
"""Monotonically increasing version number"""
version: NonNegativeInt!
"""Whether the course is visible to students"""
isPublished: Boolean!
"""Tags for categorization and search"""
tags: [String!]!
createdAt: DateTime!
updatedAt: DateTime!
```

# —— Relationships ——

```
"""The user who created this course"""
creator: User!
"""The tenant this course belongs to"""
tenant: Tenant!
"""If this course was forked, the original course"""
forkedFrom: Course
"""Ordered list of modules in this course"""
modules(
  first: PositiveInt = 50
  after: Cursor
  orderBy: ModuleOrderByInput
): ModuleConnection!
}
```

#————

# Module

#————

"""

A subdivision within a course. Modules contain media assets  
and are ordered by orderIndex.

"""

```
type Module @key(fields: "id") {
  id: UUID!
  title: String!
  description: String
  """Position within the course (0-indexed)"""
  orderIndex: NonNegativeInt!
  createdAt: DateTime!
  updatedAt: DateTime!
```

# —— Relationships ——

```
course: Course!
```

```
"""Media assets attached to this module"""
```

```
mediaAssets(  
    first: PositiveInt = 20  
    after: Cursor  
    filter: MediaAssetFilterInput  
)  
: MediaAssetConnection!  
}  
  
# -----  
# MediaAsset  
# -----
```

"""

A media file (video, audio, PDF, image, document) uploaded to the system.

Videos are transcoded to HLS; audio/video files are transcribed.

"""

```
type MediaAsset @key(fields: "id") {  
    id: UUID!  
    type: MediaType!  
    title: String!  
    description: String  
    """Original filename as uploaded"""  
    originalFilename: String  
    """MIME type of the original file"""  
    mimeType: String  
    """File size in bytes"""  
    fileSizeBytes: NonNegativeInt  
    """Duration in seconds (video/audio only)"""  
    durationSeconds: Float  
    """Current transcription status"""  
    transcriptionStatus: TranscriptionStatus  
    """Presigned URL to the HLS manifest for video playback"""  
    hlsManifestUrl: URL  
    """Presigned URL to the thumbnail image"""  
    thumbnailUrl: URL  
    """Presigned URL for direct download"""  
    downloadUrl: URL  
    """Additional technical metadata"""  
    metadata: MediaMetadata  
    createdAt: DateTime!  
    updatedAt: DateTime!
```

# ----- Relationships -----

```
"""The user who uploaded this asset"""  
uploader: User  
"""The module this asset belongs to"""  
module: Module  
"""Transcription of this asset (if available)"""
```

```
transcript: Transcript
    """Time-aligned transcript segments"""
segments(
    first: PositiveInt = 50
    after: Cursor
    filter: SegmentFilterInput
): TranscriptSegmentConnection!
}
```

```
type MediaMetadata {
    width: NonNegativeInt
    height: NonNegativeInt
    codec: String
    bitrate: NonNegativeInt
    language: String
}
```

```
# -----
# Transcript
# -----
```

"""

Full transcription of a media asset.  
Generated by Whisper or other speech-to-text models.

"""

```
type Transcript @key(fields: "id") {
    id: UUID!
    """The complete transcription text"""
    fullText: String!
    """ISO 639-1 language code (e.g., "he", "en")"""
    language: String!
    """Model used for transcription (e.g., "whisper-large-v3")"""
    modelUsed: String
    """Average confidence score (0-1)"""
    confidence: UnitFloat
    """Total word count"""
    wordCount: NonNegativeInt
    createdAt: DateTime!
```

# ----- Relationships -----

```
asset: MediaAsset!
    """Time-aligned segments of this transcript"""
segments(
    first: PositiveInt = 50
    after: Cursor
): TranscriptSegmentConnection!
}
```

```
#———  
# TranscriptSegment  
#———
```

"""

A time-aligned segment of a transcript.

Used for video synchronization, search, and knowledge graph anchoring.

"""

```
type TranscriptSegment @key(fields: "id") {  
    id: UUID!  
    """Start time in seconds from the beginning of the media"""  
    startTime: Float!  
    """End time in seconds from the beginning of the media"""  
    endTime: Float!  
    """Transcribed text for this segment"""  
    text: String!  
    """Speaker name (if speaker diarization was performed)"""  
    speaker: String  
    """Confidence score for this segment (0-1)"""  
    confidence: UnitFloat  
    createdAt: DateTime!
```

```
#——— Relationships ——  
transcript: Transcript!  
asset: MediaAsset!  
}
```

```
#———  
# Queries  
#———
```

```
type Query {  
    """Fetch a single course by ID"""  
    course(id: UUID!): Course @authenticated
```

"""

List courses within the current tenant.

Students see only published courses; instructors see all.

"""

```
courses(
```

```
    first: PositiveInt = 20
```

```
    after: Cursor
```

```
    last: PositiveInt
```

```
    before: Cursor
```

```
    filter: CourseFilterInput
```

```
    orderBy: CourseOrderByInput
```

```
): CourseConnection! @authenticated
```

```
"""Fetch a single module by ID"""
module(id: UUID!): Module @authenticated
```

```
"""Fetch a single media asset by ID"""
mediaAsset(id: UUID!): MediaAsset @authenticated
```

```
"""

```

```
List media assets within the current tenant.
```

```
Supports filtering by type and transcription status.
```

```
"""

```

```
mediaAssets(

```

```
    first: PositiveInt = 20

```

```
    after: Cursor

```

```
    filter: MediaAssetFilterInput

```

```
    orderBy: MediaAssetOrderByInput

```

```
): MediaAssetConnection! @authenticated
```

```
"""Fetch transcript segments for a time range within a media asset"""
segmentsForTimeRange(

```

```
    assetId: UUID!

```

```
    startTime: Float!

```

```
    endTime: Float!

```

```
): [TranscriptSegment!]! @authenticated
```

```
"""Full-text search across all transcripts in the current tenant"""
searchTranscripts(

```

```
    query: String!

```

```
    first: PositiveInt = 10

```

```
    after: Cursor

```

```
    assetIds: [UUID!]

```

```
): TranscriptSearchResultConnection! @authenticated
}
```

---

```
# Mutations
```

---

```
type Mutation {
```

```
# Courses
```

```
"""Create a new course. Requires INSTRUCTOR or ORG_ADMIN role."""
createCourse(input: CreateCourseInput!): Course!
```

```
    @authenticated @requiresScopes(scopes: ["course:write"])

```

```
"""Update a course. Only the creator or ORG_ADMIN can update."""

```

```
updateCourse(id: UUID!, input: UpdateCourseInput!): Course!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Soft-delete a course. Only the creator or ORG\_ADMIN can delete."""

```
deleteCourse(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Publish or unpublish a course."""

```
toggleCoursePublished(id: UUID!, isPublished: Boolean!): Course!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Fork a course (creates a copy linked to the original)."""

```
forkCourse(courseId: UUID!): Course!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

# ----- Modules -----

"""Add a module to a course."""

```
createModule(input: CreateModuleInput!): Module!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Update a module."""

```
updateModule(id: UUID!, input: UpdateModuleInput!): Module!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Delete a module (cascades to media assets)."""

```
deleteModule(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

"""Reorder modules within a course."""

```
reorderModules(courseId: UUID!, moduleOrder: [UUID!]!): [Module!]!
  @authenticated @requiresScopes(scopes: [["course:write"]])
```

# ----- Media Assets -----

"""

Initiate a media upload. Returns a presigned S3 upload URL.  
The client uploads directly to S3, then calls completeMediaUpload.

"""

```
initiateMediaUpload(input: InitiateUploadInput!): UploadTicket!
  @authenticated @requiresScopes(scopes: [["media:upload"]])
  @rateLimit(max: 20, window: "1h")
```

"""

Mark an upload as complete. Triggers transcoding/transcription pipelines.

"""

```
completeMediaUpload(uploadId: UUID!): MediaAsset!
```

```
@authenticated @requiresScopes(scopes: [["media:upload"]])\n\n\n"""Update media asset metadata.""""\nupdateMediaAsset(id: UUID!, input: UpdateMediaAssetInput!): MediaAsset!\n    @authenticated @requiresScopes(scopes: [["course:write"]])\n\n\n"""Soft-delete a media asset.""""\ndeleteMediaAsset(id: UUID!): Boolean!\n    @authenticated @requiresScopes(scopes: [["course:write"]])\n\n\n"""Re-trigger transcription for a media asset.""""\ntriggerTranscription(assetId: UUID!): MediaAsset!\n    @authenticated @requiresScopes(scopes: [["course:write"]])\n}\n\n#\n# Input Types\n#\n\ninput CreateCourseInput {\n    title: String!\n    description: String\n    syllabusJson: JSON\n    tags: [String!]\n    isPublished: Boolean = false\n}\n\ninput UpdateCourseInput {\n    title: String\n    description: String\n    syllabusJson: JSON\n    tags: [String!]\n    """Used for optimistic locking — must match current version"""\n    expectedVersion: NonNegativeInt\n}\n\ninput CreateModuleInput {\n    courseId: UUID!\n    title: String!\n    description: String\n    orderIndex: NonNegativeInt\n}\n\ninput UpdateModuleInput {\n    title: String\n    description: String\n    orderIndex: NonNegativeInt
```

```
}
```

```
input InitiateUploadInput {  
  moduleId: UUID!  
  type: MediaType!  
  title: String!  
  description: String  
  originalFilename: String!  
  mimeType: String!  
  fileSizeBytes: NonNegativeInt!  
}
```

```
input UpdateMediaAssetInput {  
  title: String  
  description: String  
}
```

```
input CourseFilterInput {  
  """Filter by published status"""  
  isPublished: Boolean  
  """Filter by creator ID"""  
  creatorId: UUID  
  """Full-text search on title and description"""  
  search: String  
  """Filter by tags (ANY match)"""  
  tags: [String!]  
}
```

```
input CourseOrderByInput {  
  field: CourseSortField!  
  direction: SortDirection!  
}
```

```
enum CourseSortField {  
  TITLE  
  CREATED_AT  
  UPDATED_AT  
  VERSION  
}
```

```
input MediaAssetFilterInput {  
  """Filter by media type"""  
  type: MediaType  
  """Filter by transcription status"""  
  transcriptionStatus: TranscriptionStatus  
  """Filter by module"""  
  moduleId: UUID
```

```
"""Full-text search on title"""
search: String
```

```
}
```

```
input MediaAssetOrderByInput {
```

```
    field: MediaAssetSortField!
    direction: SortDirection!
```

```
}
```

```
enum MediaAssetSortField {
```

```
    TITLE
```

```
    CREATED_AT
```

```
    FILE_SIZE
```

```
    DURATION
```

```
    TYPE
```

```
}
```

```
input ModuleOrderByInput {
```

```
    field: ModuleSortField!
```

```
    direction: SortDirection!
```

```
}
```

```
enum ModuleSortField {
```

```
    ORDER_INDEX
```

```
    TITLE
```

```
    CREATED_AT
```

```
}
```

```
input SegmentFilterInput {
```

```
    """Filter by speaker name"""
speaker: String
```

```
    """Full-text search within segment text"""
search: String
```

```
    """Minimum confidence threshold"""
minConfidence: UnitFloat
```

```
}
```

---

```
# Connection Types
```

---

```
type CourseConnection {
```

```
    edges: [CourseEdge!]!
```

```
    pageInfo: PageInfo!
```

```
    totalCount: NonNegativeInt
```

```
}
```

```
type CourseEdge {  
    node: Course!  
    cursor: Cursor!  
}
```

```
type ModuleConnection {  
    edges: [ModuleEdge!]!  
    pageInfo: PageInfo!  
    totalCount: NonNegativeInt  
}
```

```
type ModuleEdge {  
    node: Module!  
    cursor: Cursor!  
}
```

```
type MediaAssetConnection {  
    edges: [MediaAssetEdge!]!  
    pageInfo: PageInfo!  
    totalCount: NonNegativeInt  
}
```

```
type MediaAssetEdge {  
    node: MediaAsset!  
    cursor: Cursor!  
}
```

```
type TranscriptSegmentConnection {  
    edges: [TranscriptSegmentEdge!]!  
    pageInfo: PageInfo!  
    totalCount: NonNegativeInt  
}
```

```
type TranscriptSegmentEdge {  
    node: TranscriptSegment!  
    cursor: Cursor!  
}
```

```
# ----- Upload Ticket -----
```

```
"""
```

Returned by `initiateMediaUpload`. Contains a presigned S3 URL for the client to upload the file directly.

```
"""
```

```
type UploadTicket {  
    """Unique ID for this upload session"""  
    uploadId: UUID!
```

```

"""Presigned PUT URL for direct upload to S3"""
presignedUrl: URL!
"""Expiration time of the presigned URL"""
expiresAt: DateTime!
"""Required headers to include in the PUT request"""
requiredHeaders: JSON
}

```

# ----- Transcript Search -----

```

type TranscriptSearchResult {
  segment: TranscriptSegment!
  asset: MediaAsset!
  """Relevance score"""
  score: Float!
  """Highlighted text snippet with <mark> tags"""
  highlightedText: String!
}

```

```

type TranscriptSearchResultConnection {
  edges: [TranscriptSearchResultEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

```

```

type TranscriptSearchResultEdge {
  node: TranscriptSearchResult!
  cursor: Cursor!
}

```

## 9. Subgraph: Annotation (Layers, Sketches, Comments)

**Owns:** `Annotation` **Port:** 4003 **Database tables:** `annotations`

### Schema Definition

```
graphql
```

```
# ----- File: apps/subgraph-annotation/src/schema.graphql -----
```

```
extend schema
```

```
@link(url: "https://specs.apollo.dev/federation/v2.7",
import: ["@key", "@external", "@requires", "@authenticated", "@requiresScopes"])
```

```
# -----
```

```
# Entity Stubs
```

```
# -----
```

```
type User @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
type MediaAsset @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
# -----
```

```
# Extend entities with annotation fields
```

```
# -----
```

```
extend type User {
  """All annotations created by this user"""
  annotations(
    first: PositiveInt = 20
    after: Cursor
    filter: AnnotationFilterInput
  ): AnnotationConnection!
}
```

```
extend type MediaAsset {
  """All annotations on this media asset"""
  annotations(
    first: PositiveInt = 20
    after: Cursor
    filter: AnnotationFilterInput
  ): AnnotationConnection!
```

```
"""Annotations at a specific timestamp (±tolerance seconds)"""
annotationsAtTimestamp(

```

```
  timestamp: Float!
  toleranceSeconds: Float = 2.0
  layers: [AnnotationLayer!]!
): [Annotation!]!
```

```
}
```

```
#  
# Annotation  
#
```

"""

An annotation is any user or AI marking on content:

text notes, sketches, bookmarks, spatial comments, or links.

Annotations are layered (personal, shared, instructor, AI-generated)

and can be temporally/spatially anchored to media.

"""

```
type Annotation @key(fields: "id") {
```

```
  id: UUID!
```

```
  type: AnnotationType!
```

```
  """Visibility layer"""
  layer: AnnotationLayer!
```

# —— Content ——

```
  """Text content or markdown (for text, link, bookmark types)"""
  content: String
```

```
  """Konva.js sketch data (for sketch type)"""
  sketchData: JSON
```

# —— Temporal anchoring ——

```
  """Start time in seconds (null for non-temporal annotations)"""
  timestampStart: Float
```

```
  """End time in seconds (null for point annotations)"""
  timestampEnd: Float
```

# —— Spatial anchoring (video frame) ——

```
  """Normalized X position (0-1) on the video frame"""
  spatialX: Float
```

```
  """Normalized Y position (0-1) on the video frame"""
  spatialY: Float
```

# —— PDF anchoring ——

```
  """Page number (1-indexed) for PDF annotations"""
  pageNumber: NonNegativeInt
```

```
  """Character offset start for text range selection"""
  textRangeStart: NonNegativeInt
```

```
  """Character offset end for text range selection"""
  textRangeEnd: NonNegativeInt
```

# —— Metadata ——

```
  metadata: AnnotationMetadata
```

```
  createdAt: DateTime!
```

```
  updatedAt: DateTime!
```

```
# ----- Relationships -----  
"""The user who created this annotation"""  
author: User!  
"""The media asset this annotation is attached to"""  
asset: MediaAsset!  
"""Parent annotation (for reply chains)"""  
parent: Annotation  
"""Replies to this annotation"""  
replies(  
    first: PositiveInt = 20  
    after: Cursor  
) : AnnotationConnection!  
"""Agent that created this annotation (if AI-generated)"""  
agentId: UUID  
}
```

```
type AnnotationMetadata {  
    color: String  
    pinned: Boolean  
    resolvedAt: DateTime  
    linkedConceptIds: [UUID!]  
    aiConfidence: UnitFloat  
}
```

```
# -----  
# Queries  
# -----
```

```
type Query {  
    """Fetch a single annotation by ID"""  
    annotation(id: UUID!): Annotation @authenticated
```

.....

List annotations with filtering.

Students see their own personal layer + shared/instructor/AI layers.

.....

```
annotations(  
    first: PositiveInt = 20  
    after: Cursor  
    last: PositiveInt  
    before: Cursor  
    filter: AnnotationFilterInput  
    orderBy: AnnotationOrderByInput  
) : AnnotationConnection! @authenticated
```

```
"""Get annotation thread (all replies to a root annotation)"""
```

```
annotationThread(rootId: UUID!): [Annotation!]! @authenticated
}
```

```
# -----
# Mutations
# -----
```

```
type Mutation {
    """Create a new annotation on a media asset."""
    createAnnotation(input: CreateAnnotationInput!): Annotation!
        @authenticated
```

```
"""
Update an annotation. Users can only update their own annotations.
Instructors/admins can update shared layer annotations.
```

```
"""
updateAnnotation(id: UUID!, input: UpdateAnnotationInput!): Annotation!
    @authenticated
```

```
"""
Delete an annotation and all its replies (cascade).
Users can only delete their own annotations.
```

```
"""
deleteAnnotation(id: UUID!): Boolean!
    @authenticated
```

```
"""
Pin or unpin an annotation (instructors/admins)."""
toggleAnnotationPin(id: UUID!, pinned: Boolean!): Annotation!
    @authenticated @requiresScopes(scopes: [[ "annotation:write" ]])
```

```
"""
Mark an annotation as resolved."""
resolveAnnotation(id: UUID!): Annotation!
    @authenticated
```

```
"""
Batch-move annotations between layers (instructor/admin only)."""
moveAnnotationsToLayer(
    annotationIds: [UUID!]!
    targetLayer: AnnotationLayer!
): [Annotation!]!
    @authenticated @requiresScopes(scopes: [[ "annotation:write" ]])
```

```
}
```

---

```
# -----
# Input Types
# -----
```

```
input CreateAnnotationInput {
```

```
assetId: UUID!
type: AnnotationType!
layer: AnnotationLayer = PERSONAL
content: String
sketchData: JSON
timestampStart: Float
timestampEnd: Float
spatialX: Float
spatialY: Float
pageNumber: NonNegativeInt
textRangeStart: NonNegativeInt
textRangeEnd: NonNegativeInt
parentId: UUID
metadata: AnnotationMetadataInput
}
```

```
input UpdateAnnotationInput {
  content: String
  sketchData: JSON
  timestampStart: Float
  timestampEnd: Float
  spatialX: Float
  spatialY: Float
  metadata: AnnotationMetadataInput
}
```

```
input AnnotationMetadataInput {
  color: String
  linkedConceptIds: [UUID!]
}
```

```
input AnnotationFilterInput {
  """Filter by media asset"""
  assetId: UUID
  """Filter by annotation type"""
  type: AnnotationType
  """Filter by layer(s)"""
  layers: [AnnotationLayer!]
  """Filter by author"""
  userId: UUID
  """Only show pinned annotations"""
  pinnedOnly: Boolean
  """Only show unresolved annotations"""
  unresolvedOnly: Boolean
  """Only show root annotations (no replies)"""
  rootOnly: Boolean
  """Full-text search in content"""
}
```

```
search: String
    """Filter by time range (start)"""
timestampFrom: Float
    """Filter by time range (end)"""
timestampTo: Float
    """Filter by page number (PDF)"""
pageNumber: NonNegativeInt
}
```

```
input AnnotationOrderByInput {
    field: AnnotationSortField!
    direction: SortDirection!
}
```

```
enum AnnotationSortField {
    CREATED_AT
    TIMESTAMP_START
    PAGE_NUMBER
    LAYER
}
```

```
# _____
# Connection Types
# _____
```

```
type AnnotationConnection {
    edges: [AnnotationEdge!]!
    pageInfo: PageInfo!
    totalCount: NonNegativeInt
}
```

```
type AnnotationEdge {
    node: Annotation!
    cursor: Cursor!
}
```

## 10. Subgraph: Collaboration (CRDT, Real-time)

**Owns:** `CollabDocument`, `CollabSession` **Port:** 4004 **Database tables:** `collab_documents`, `crdt_updates`, `collab_sessions`

### Schema Definition

```
graphql
```

```
# ----- File: apps/subgraph-collaboration/src/schema.graphql -----
```

```
extend schema
```

```
@link(url: "https://specs.apollo.dev/federation/v2.7",
  import: ["@key", "@external", "@authenticated", "@requiresScopes"])
```

```
# -----
```

```
# Entity Stubs
```

```
# -----
```

```
type User @key(fields: "id", resolvable: false) {
```

```
  id: UUID!
```

```
}
```

```
# -----
```

```
# CollabDocument
```

```
# -----
```

```
"""
```

A collaborative document backed by Yjs CRDT.

Used for shared annotation layers, course notes, and real-time editing.

```
"""
```

```
type CollabDocument @key(fields: "id") {
```

```
  id: UUID!
```

```
    """Hocuspocus document name (unique within tenant)"""
  documentName: String!
```

```
    """Human-readable title"""
  title: String
```

```
    """Type of entity this document is associated with"""
  entityType: String
```

```
    """ID of the associated entity"""
  entityId: UUID
```

```
    """Number of pending (uncompacted) CRDT updates"""
  pendingUpdates: NonNegativeInt!
```

```
    """Whether a compacted state vector snapshot exists"""
  hasSnapshot: Boolean!
```

```
  createdAt: DateTime!
```

```
  updatedAt: DateTime!
```

```
# ----- Relationships -----
```

```
    """Currently active collaboration sessions"""
  activeSessions: [CollabSession!]!
```

```
    """Number of active collaborators"""
  activeCollaboratorCount: NonNegativeInt!
```

```
}
```

```
#-----  
# CollabSession  
#-----  
  
"""  
Represents an active user session within a collaborative document.  
Ephemeral — cleaned up when the user disconnects.  
"""  
  
type CollabSession @key(fields: "id") {  
    id: UUID!  
    """The user participating in this session"""  
    user: User!  
    """The document being collaborated on"""  
    document: CollabDocument!  
    """Current cursor position of this collaborator"""  
    cursorPosition: CursorPosition  
    """When the user connected"""  
    connectedAt: DateTime!  
    """Last heartbeat from the user"""  
    lastSeenAt: DateTime!  
}  
  
type CursorPosition {  
    x: Float  
    y: Float  
    """Video timestamp if collaborating on video annotations"""  
    timestamp: Float  
    """PDF page number if collaborating on PDF annotations"""  
    pageNumber: NonNegativeInt  
    """Display name for the cursor label"""  
    userName: String  
    """Color assigned to this collaborator"""  
    color: String  
}  
  
#-----  
# Queries  
#-----  
  
type Query {  
    """Fetch a collaborative document by ID"""  
    collabDocument(id: UUID!): CollabDocument @authenticated  
  
    """Fetch a collaborative document by its document name"""  
    collabDocumentByName(documentName: String!): CollabDocument @authenticated  
  
    """List collaborative documents for an entity"""
```

```

collabDocumentsForEntity(
    entityType: String!
    entityId: UUID!
): [CollabDocument!]! @authenticated

"""Get the Hocuspocus WebSocket URL for a document"""
collabConnectionInfo(
    documentId: UUID!
): CollabConnectionInfo! @authenticated
}

type CollabConnectionInfo {
    """WebSocket URL for Hocuspocus connection"""
    wsUrl: URL!
    """Authentication token for the WebSocket connection"""
    authToken: String!
    """Document name to use in the Hocuspocus protocol"""
    documentName: String!
}

# -----
# Mutations
# -----



type Mutation {
    """Create a new collaborative document"""
    createCollabDocument(input: CreateCollabDocumentInput!): CollabDocument!
        @authenticated

    """Force-compact CRDT updates into a snapshot (admin)"""
    compactCollabDocument(documentId: UUID!): CollabDocument!
        @authenticated @requiresScopes(scopes: [[ "org:manage" ]])
}

input CreateCollabDocumentInput {
    documentName: String!
    title: String
    entityType: String
    entityId: UUID
}

```

## 11. Subgraph: Agent (Definitions, Executions, Templates)

**Owns:** [AgentDefinition](#), [AgentExecution](#) **Port:** 4005 **Database tables:** [agent\\_definitions](#), [agent\\_executions](#)

## Schema Definition

graphql

```
# ----- File: apps/subgraph-agent/src/schema.graphql -----
```

```
extend schema
```

```
  @link(url: "https://specs.apollo.dev/federation/v2.7",
    import: ["@key", "@external", "@authenticated", "@requiresScopes"])
```

```
# -----
```

```
# Entity Stubs
```

```
# -----
```

```
type User @key(fields: "id", resolvable: false) {
```

```
  id: UUID!
```

```
}
```

```
type Course @key(fields: "id", resolvable: false) {
```

```
  id: UUID!
```

```
}
```

```
# -----
```

```
# Extend User with agent fields
```

```
# -----
```

```
extend type User {
```

```
    """Agent definitions created by this user"""
    agentDefinitions(
```

```
      first: PositiveInt = 20
```

```
      after: Cursor
```

```
      filter: AgentDefinitionFilterInput
```

```
    ): AgentDefinitionConnection!
```

```
    """Agent execution history for this user"""
    agentExecutions(
```

```
      first: PositiveInt = 20
```

```
      after: Cursor
```

```
      filter: AgentExecutionFilterInput
```

```
    ): AgentExecutionConnection!
```

```
}
```

```
# -----
```

```
# AgentDefinition
```

```
# -----
```

```
"""
```

A user-created or built-in AI agent definition.

Agents process content using LLMs based on configurable templates  
(Chavruta debate partner, summarizer, quiz master, etc.).

```
"""
type AgentDefinition @key(fields: "id") {
    id: UUID!
    name: String!
    description: String
    """Base template type"""
    templateType: AgentTemplate!
    """Whether this is a system-provided built-in agent"""
    isBuiltIn: Boolean!
    """Whether this agent is visible to all users in the tenant"""
    isPublic: Boolean!

    # ----- Configuration -----
    config: AgentConfig!

    # ----- Usage Stats -----
    """Total number of times this agent has been executed"""
    executionCount: NonNegativeInt!
    """Average execution time in milliseconds"""
    avgExecutionMs: Float
    """Last time this agent was executed"""
    lastExecutedAt: DateTime

    createdAt: DateTime!
    updatedAt: DateTime!

    # ----- Relationships -----
    """User who created this agent definition"""
    creator: User!
    """Recent executions of this agent"""
    recentExecutions(
        first: PositiveInt = 10
        after: Cursor
    ): AgentExecutionConnection!
}

type AgentConfig {
    """Scope of data the agent can access"""
    dataScope: AgentDataScope!
    """Custom Cypher query for CUSTOM_QUERY scope"""
    customQuery: String
    """What triggers the agent"""
    triggerType: AgentTrigger!
    """How the agent's output is displayed"""
    outputFormat: AgentOutputFormat!
    """Agent personality description"""
    personality: String
}
```

```
"""Difficulty level for educational agents"""
difficulty: AgentDifficulty
"""Override the default LLM model"""
modelOverride: String
"""Maximum tokens for the response"""
maxTokens: NonNegativeInt
"""LLM temperature (0.0-2.0)"""
temperature: Float
"""Custom system prompt"""
systemPrompt: String
"""List of MCP tools this agent can use"""
toolsEnabled: [String!]
}
```

```
enum AgentDataScope {
    MY_NOTES
    COURSE
    ORGANIZATION
    CUSTOM_QUERY
}
```

```
enum AgentDifficulty {
    BEGINNER
    INTERMEDIATE
    ADVANCED
}
```

```
# -----
# AgentExecution
# -----
```

```
"""
A record of a single agent execution.
Contains input, output, performance metrics, and debugging info.
"""


```

```
type AgentExecution @key(fields: "id") {
    id: UUID!
    """Current status of the execution"""
    status: ExecutionStatus!

    # ----- Input / Output -----
    """Input provided to the agent"""
    input: JSON!
    """Agent's output (null while running)"""
    output: JSON
    """Chain-of-thought reasoning (for debugging)"""
    reasoning: String
}
```

```
"""Error message if execution failed"""
errorMessage: String

# ----- Performance -----
"""Number of tokens consumed"""
tokensUsed: NonNegativeInt
"""Execution time in milliseconds"""
executionMs: NonNegativeInt
"""LLM model used for this execution"""
modelUsed: String

createdAt: DateTime!
updatedAt: DateTime!

# ----- Relationships -----
"""The agent definition that was executed"""
agent: AgentDefinition!
"""The user who triggered the execution"""
user: User!
}

# -----
# Queries
# -----



type Query {
    """Fetch an agent definition by ID"""
    agentDefinition(id: UUID!): AgentDefinition @authenticated

    """List available agent definitions (includes public + own)"""
    agentDefinitions(
        first: PositiveInt = 20
        after: Cursor
        filter: AgentDefinitionFilterInput
        orderBy: AgentDefinitionOrderByInput
    ): AgentDefinitionConnection! @authenticated

    """List built-in agent templates"""
    agentTemplates: [AgentTemplateInfo!]! @authenticated

    """Fetch a specific agent execution"""
    agentExecution(id: UUID!): AgentExecution @authenticated

    """List agent executions with filtering"""
    agentExecutions(
        first: PositiveInt = 20
        after: Cursor
    ): AgentExecutionConnection! @authenticated
}
```

```

filter: AgentExecutionFilterInput
): AgentExecutionConnection! @authenticated
}

"""Information about a built-in agent template"""
type AgentTemplateInfo {
  templateType: AgentTemplate!
  name: String!
  description: String!
  """Default configuration for this template"""
  defaultConfig: AgentConfig!
  """Icon identifier for UI"""
  icon: String!
}

# -----
# Mutations
# -----



type Mutation {
  """Create a new agent definition"""
  createAgentDefinition(input: CreateAgentDefinitionInput!): AgentDefinition!
    @authenticated @requiresScopes(scopes: [[ "agent:write" ]])

  """Update an agent definition (creator only)"""
  updateAgentDefinition(id: UUID!, input: UpdateAgentDefinitionInput!): AgentDefinition!
    @authenticated @requiresScopes(scopes: [[ "agent:write" ]])

  """Delete an agent definition (creator only, cascades executions)"""
  deleteAgentDefinition(id: UUID!): Boolean!
    @authenticated @requiresScopes(scopes: [[ "agent:write" ]])

  """
  Execute an agent. Returns immediately with QUEUED status.
  Use the agentExecutionUpdated subscription to track progress.
  """

  executeAgent(input: ExecuteAgentInput!): AgentExecution!
    @authenticated @requiresScopes(scopes: [[ "agent:execute" ]])
    @rateLimit(max: 30, window: "1h")

  """Cancel a running agent execution"""
  cancelAgentExecution(executionId: UUID!): AgentExecution!
    @authenticated
}

# -----
# Input Types

```

#

---

```
input CreateAgentDefinitionInput {
    name: String!
    description: String
    templateType: AgentTemplate!
    isPublic: Boolean = false
    config: AgentConfigInput!
}

input UpdateAgentDefinitionInput {
    name: String
    description: String
    isPublic: Boolean
    config: AgentConfigInput
}

input AgentConfigInput {
    dataScope: AgentDataScope!
    customQuery: String
    triggerType: AgentTrigger!
    outputFormat: AgentOutputFormat!
    personality: String
    difficulty: AgentDifficulty
    modelOverride: String
    maxTokens: NonNegativeInt
    temperature: Float
    systemPrompt: String
    toolsEnabled: [String!]
}

input ExecuteAgentInput {
    agentId: UUID!
    """Context-specific input (e.g., annotation ID, course ID, free-text prompt)"""
    input: JSON!
}

input AgentDefinitionFilterInput {
    templateType: AgentTemplate
    isBuiltIn: Boolean
    isPublic: Boolean
    creatorId: UUID
    search: String
}

input AgentDefinitionOrderByInput {
    field: AgentDefinitionSortField!
```

```

direction: SortDirection!
}

enum AgentDefinitionSortField {
  NAME
  CREATED_AT
  EXECUTION_COUNT
  TEMPLATE_TYPE
}

input AgentExecutionFilterInput {
  agentId: UUID
  status: ExecutionStatus
  userId: UUID
}

#-----
# Connection Types
#-----


type AgentDefinitionConnection {
  edges: [AgentDefinitionEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

type AgentDefinitionEdge {
  node: AgentDefinition!
  cursor: Cursor!
}

type AgentExecutionConnection {
  edges: [AgentExecutionEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

type AgentExecutionEdge {
  node: AgentExecution!
  cursor: Cursor!
}

```

## 12. Subgraph: Knowledge (Graph, Embeddings, Semantic Search)

**Owns:** (Concept), (Person), (Term), (Source), (TopicCluster), (KnowledgeRelation), (SemanticSearchResult) **Port:** 4006

**Database:** Apache AGE graph + pgvector embeddings

## Schema Definition

```
graphql
```

```
# ----- File: apps/subgraph-knowledge/src/schema.graphql -----
```

```
extend schema
```

```
  @link(url: "https://specs.apollo.dev/federation/v2.7",
    import: ["@key", "@external", "@requires", "@authenticated", "@requiresScopes"])
```

```
#-----
```

```
# Entity Stubs
```

```
#-----
```

```
type MediaAsset @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
type TranscriptSegment @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
type Annotation @key(fields: "id", resolvable: false) {
  id: UUID!
}
```

```
#-----
```

```
# Extend MediaAsset with knowledge fields
```

```
#-----
```

```
extend type MediaAsset {
  """Concepts mentioned in this media asset's transcript"""
  mentionedConcepts(
    first: PositiveInt = 20
    after: Cursor
  ): ConceptConnection!
```

```
"""
```

```
Semantic context for a specific timestamp.
```

```
Powers the "Contextual Copilot" sidebar.
```

```
"""
```

```
semanticContextAt(timestamp: Float!): SemanticContext!
}
```

```
extend type TranscriptSegment {
```

```
  """Concepts mentioned in this specific segment"""
  mentionedConcepts: [ConceptMention!]!
```

```
}
```

```
extend type Annotation {
```

"""Concepts linked to this annotation"""

linkedConcepts: [Concept!]!

}

# —————

# Concept (Knowledge Graph Vertex)

# —————

"""

A semantic concept in the knowledge graph.

Concepts are extracted from content (AI or manual) and form the backbone of the knowledge graph.

"""

```
type Concept @key(fields: "id") {
    id: UUID!
    """Primary label (Hebrew/English)"""
    label: String!
    """Detailed description of the concept"""
    description: String
    """Alternative names, abbreviations, synonyms"""
    aliases: [String!]!
    """Academic or knowledge domain (e.g., "Physics", "Talmud")"""
    domain: String
    """Confidence: 1.0 = manually created, <1.0 = AI-extracted"""
    confidence: UnitFloat!
    """How this concept was created"""
    sourceType: ConceptSourceType!
    createdAt: DateTime!
```

# ————— Knowledge Graph Relationships —————

"""Concepts related to this one (with strength scores)"""

relatedConcepts(

```
    first: PositiveInt = 20
    after: Cursor
    minStrength: UnitFloat
): KnowledgeRelationConnection!
```

"""Concepts that contradict this one"""

contradictions: [Contradiction!]!

"""Prerequisites for understanding this concept"""

prerequisites: [PrerequisiteLink!]!

"""Concepts that require this concept as a prerequisite"""

dependents: [PrerequisiteLink!]!

"""Full prerequisite learning path (up to 5 levels deep)"""

```
prerequisiteChain: [Concept!]!
```

```
"""Media segments where this concept is mentioned"""
mentions(
```

```
    first: PositiveInt = 20
```

```
    after: Cursor
```

```
): ConceptMentionConnection!
```

```
"""Topic cluster this concept belongs to"""
topicCluster: TopicCluster
```

```
}
```

```
enum ConceptSourceType {
```

```
MANUAL
```

```
AI_EXTRACTED
```

```
IMPORTED
```

```
}
```

---

```
#-----
```

```
# Knowledge Graph Edge Types
```

---

```
#-----
```

```
"""A relationship between two concepts"""
type KnowledgeRelation {
```

```
    """The related concept"""
    concept: Concept!
```

```
    """Semantic similarity or relationship strength (0-1)"""
    strength: UnitFloat!
```

```
    """How the relationship was established"""
    method: RelationMethod!
```

```
    """Human-readable explanation"""
    evidence: String
```

```
}
```

```
enum RelationMethod {
```

```
MANUAL
```

```
EMBEDDING_SIMILARITY
```

```
CO_OCCURRENCE
```

```
CITATION
```

```
}
```

```
"""A contradiction between two concepts"""
type Contradiction {
```

```
    """The contradicting concept"""
    concept: Concept!
```

```
    """Explanation of why they contradict"""
    evidence: String!
```

```
"""Severity of the contradiction"""
severity: ContradictionSeverity!
"""How the contradiction was detected"""
detectedBy: ContradictionDetector!
"""Source media asset A"""
sourceAssetA: MediaAsset
"""Source media asset B"""
sourceAssetB: MediaAsset
}
```

```
enum ContradictionSeverity {
    MINOR
    MODERATE
    MAJOR
}
```

```
enum ContradictionDetector {
    MANUAL
    CHAVRUTA_AGENT
    AI_INFERENCE
}
```

```
"""A prerequisite relationship between concepts"""
type PrerequisiteLink {
    """The prerequisite (or dependent) concept"""
    concept: Concept!
    """How strongly this is recommended"""
    strength: PrerequisiteStrength!
}
```

```
enum PrerequisiteStrength {
    REQUIRED
    RECOMMENDED
    OPTIONAL
}
```

```
"""A mention of a concept in a specific transcript segment"""
type ConceptMention {
    concept: Concept!
    segment: TranscriptSegment!
    asset: MediaAsset!
    startTime: Float!
    endTime: Float!
    confidence: UnitFloat!
}
```

```
"""Auto-generated grouping of related concepts"""

```

```
type TopicCluster {
  id: UUID!
  label: String!
  conceptCount: NonNegativeInt!
  coherenceScore: UnitFloat!
  """Concepts in this cluster"""
  concepts(
    first: PositiveInt = 20
    after: Cursor
  ): ConceptConnection!
}
```

```
# -----
# Person (Knowledge Graph Vertex)
# -----
```

```
"""
A person referenced in content — author, speaker, historical figure.
Stored as a vertex in the Apache AGE knowledge graph.
```

```
"""
type Person @key(fields: "id") {
  id: UUID!
  """Full name of the person"""
  name: String!
  """Role in context (e.g., "Speaker", "Author", "Historical")"""
  role: String
  """External reference URL (e.g., Wikipedia)"""
  externalUrl: URL
  """Content this person authored"""
  authoredSources(
    first: PositiveInt = 10
    after: Cursor
  ): SourceConnection!
  """Concepts this person is associated with"""
  relatedConcepts: [Concept!]!
  createdAt: DateTime!
}
```

```
# -----
# Term (Knowledge Graph Vertex)
# -----
```

```
"""
A domain-specific term — narrower and more precise than a Concept.
Terms are leaf nodes in the knowledge graph, linked to parent Concepts.
```

```
"""
type Term @key(fields: "id") {
```

```
id: UUID!
"""Display label of the term"""
label: String!
"""Knowledge domain (e.g., "Physics", "Talmud")"""
domain: String!
"""Formal definition of the term"""
definition: String
"""Parent concept this term belongs to"""
parentConcept: Concept
createdAt: DateTime!
}
```

---

```
# _____
# Source (Knowledge Graph Vertex)
# _____
```

"""\nAn external reference — book, paper, URL, lecture, or internal asset.\nSources are vertices in the knowledge graph and can be linked to\nConcepts via CITES or REFERS\_TO edges.

"""\n

```
type Source @key(fields: "id") {
  id: UUID!
  """Title of the source material"""
  title: String!
  """Type of source"""
  type: SourceType!
  """External URL (for web resources)"""
  externalUrl: URL
  """Internal media asset (if type is INTERNAL_ASSET)"""
  internalAsset: MediaAsset
  """ISBN for books"""
  isbn: String
  """DOI for papers"""
  doi: String
  """Person(s) who authored this source"""
  authors: [Person!]!
  """Concepts cited or referenced in this source"""
  citedConcepts(
    first: PositiveInt = 20
    after: Cursor
  ): ConceptConnection!
  createdAt: DateTime!
}
```

```
enum SourceType {
  BOOK
```

```
PAPER  
URL  
LECTURE  
INTERNAL_ASSET  
}
```

```
#-----  
# Semantic Context (Contextual Copilot)  
#-----
```

"""

Rich semantic context for a point in time within a media asset.

Combines knowledge graph traversal with vector similarity.

Powers the "Contextual Copilot" sidebar in the video player.

"""

```
type SemanticContext {  
    """Concepts being discussed at this timestamp"""  
    activeConcepts: [ConceptWithRelations!]!  
    """Suggested related content from other assets"""  
    relatedContent: [RelatedContentSuggestion!]!  
    """Detected contradictions at this point"""  
    activeContradictions: [Contradiction!]!  
    """Suggested prerequisite concepts the viewer may need"""  
    suggestedPrerequisites: [Concept!]!  
}
```

```
type ConceptWithRelations {  
    concept: Concept!  
    """How this concept connects to other active concepts"""  
    connections: [ConceptConnection!]!  
}
```

```
type ConceptConnection {  
    relatedConcept: Concept!  
    relationType: String!  
    strength: UnitFloat!  
}
```

```
type RelatedContentSuggestion {  
    """The related media asset"""  
    asset: MediaAsset!  
    """The segment within the asset that is most relevant"""  
    segment: TranscriptSegment  
    """Why this content is suggested"""  
    reason: String!  
    """Relevance score"""  
    relevance: UnitFloat!
```

```
}
```

```
#-----  
# Semantic Search  
#-----
```

```
"""
```

```
Result of a semantic (vector) search across content.
```

```
Combines pgvector nearest-neighbor search with knowledge graph context.
```

```
"""
```

```
type SemanticSearchResult {  
    """The transcript segment that matched"""  
    segment: TranscriptSegment!  
    """The media asset containing the match"""  
    asset: MediaAsset!  
    """The text that was matched"""  
    matchedText: String!  
    """Cosine similarity score"""  
    similarity: UnitFloat!  
    """Concepts mentioned in the matched segment"""  
    relatedConcepts: [Concept!]!  
}
```

```
#-----  
# Queries  
#-----
```

```
type Query {  
    """Fetch a concept by ID"""  
    concept(id: UUID!): Concept @authenticated  
  
    """Search concepts by label, alias, or description"""  
    concepts(  
        first: PositiveInt = 20  
        after: Cursor  
        filter: ConceptFilterInput  
        orderBy: ConceptOrderByInput  
    ): ConceptConnection! @authenticated
```

```
"""
```

```
Semantic search across all content using natural language.
```

```
Uses pgvector HNSW index for approximate nearest neighbor search.
```

```
"""
```

```
semanticSearch(  
    """Natural language query (will be embedded by the server)"""  
    query: String!  
    first: PositiveInt = 10
```

```
after: Cursor
"""Minimum similarity threshold (0-1)"""
minSimilarity: UnitFloat = 0.7
"""Limit search to specific asset IDs"""
assetIds: [UUID!]
): SemanticSearchResultConnection! @authenticated
```

```
"""
HybridRAG search: combines vector search + knowledge graph traversal.
Best for complex queries that need both semantic and structural context.
```

```
"""
hybridSearch(
    query: String!
    first: PositiveInt = 10
    """How many graph hops to follow from vector results"""
    graphDepth: PositiveInt = 2
): HybridSearchResultConnection! @authenticated
```

```
"""
Find concepts related to a given concept (graph traversal)
relatedConcepts(
    conceptId: UUID!
    """Max graph traversal depth"""
    maxDepth: PositiveInt = 3
    first: PositiveInt = 50
): KnowledgeRelationConnection! @authenticated
```

```
"""
Find contradictions for a concept
contradictions(
    conceptId: UUID!
): [Contradiction!]! @authenticated
```

```
"""
Build a learning path (prerequisite chain) for a concept
learningPath(
    conceptId: UUID!
    maxDepth: PositiveInt = 5
): [Concept!]! @authenticated
```

```
"""
List topic clusters
topicClusters(
    first: PositiveInt = 20
    after: Cursor
): TopicClusterConnection! @authenticated
```

```
# —— Person queries ——
```

```
"""
Fetch a person by ID
person(id: UUID!): Person @authenticated
```

```
"""Search people referenced in the knowledge graph"""
people(
  first: PositiveInt = 20
  after: Cursor
  search: String
): PersonConnection! @authenticated

# ----- Term queries -----

"""Fetch a term by ID"""
term(id: UUID!): Term @authenticated

"""Search domain-specific terms"""
terms(
  first: PositiveInt = 20
  after: Cursor
  domain: String
  search: String
): TermConnection! @authenticated

# ----- Source queries -----

"""Fetch a source by ID"""
source(id: UUID!): Source @authenticated

"""Search sources (books, papers, URLs, etc.)"""
sources(
  first: PositiveInt = 20
  after: Cursor
  type: SourceType
  search: String
): SourceConnection! @authenticated
}

# -----
# Mutations
# -----
```

```
type Mutation {
  """Manually create a concept in the knowledge graph"""
  createConcept(input: CreateConceptInput!): Concept!
    @authenticated @requiresScopes(scopes: ["knowledge:write"])

  """Update a concept's label, description, aliases, or domain"""
  updateConcept(id: UUID!, input: UpdateConceptInput!): Concept!
    @authenticated @requiresScopes(scopes: ["knowledge:write"])
```

```
"""Delete a concept and all its edges from the knowledge graph"""
deleteConcept(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])

"""Create a relationship between two concepts"""
createRelation(input: CreateRelationInput!): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])

"""Remove a relationship between two concepts"""
deleteRelation(
  fromConceptId: UUID!
  toConceptId: UUID!
  relationType: String!
): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

```
"""Mark a contradiction between two concepts"""
createContradiction(input: CreateContradictionInput!): Contradiction!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

Trigger re-indexing of embeddings for a media asset.  
Used after transcript updates or corrections.

```
reindexAssetEmbeddings(assetId: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

```
"""Review and approve/reject an AI-inferred relationship"""
reviewInferredRelation(
  fromConceptId: UUID!
  toConceptId: UUID!
  approved: Boolean!
): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

# ----- Person mutations -----

```
"""Create a person reference in the knowledge graph"""
createPerson(input: CreatePersonInput!): Person!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

```
"""Update a person's details"""
updatePerson(id: UUID!, input: UpdatePersonInput!): Person!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

```
"""Delete a person from the knowledge graph"""
deletePerson(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [[ "knowledge:write" ]])
```

```
deletePerson(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

# ----- Term mutations -----

"""Create a domain-specific term"""

```
createTerm(input: CreateTermInput!): Term!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Update a term's details"""

```
updateTerm(id: UUID!, input: UpdateTermInput!): Term!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Delete a term from the knowledge graph"""

```
deleteTerm(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

# ----- Source mutations -----

"""Create a source reference (book, paper, URL, etc.)"""

```
createSource(input: CreateSourceInput!): Source!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Update a source's details"""

```
updateSource(id: UUID!, input: UpdateSourceInput!): Source!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Delete a source from the knowledge graph"""

```
deleteSource(id: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Link a person as author of a source (AUTHORED\_BY edge)"""

```
linkAuthorToSource(personId: UUID!, sourceId: UUID!): Boolean!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

"""Link a source to a concept (CITES edge)"""

```
linkSourceToConcept(
  sourceId: UUID!
  conceptId: UUID!
  page: Int
  timestamp: Float
  context: String
): Boolean!
  @authenticated @requiresScopes(scopes: [["knowledge:write"]])
```

}

# -----

```
# Input Types
```

```
#-----
```

```
input CreateConceptInput {
```

```
  label: String!
```

```
  description: String
```

```
  aliases: [String!]
```

```
  domain: String
```

```
}
```

```
input UpdateConceptInput {
```

```
  label: String
```

```
  description: String
```

```
  aliases: [String!]
```

```
  domain: String
```

```
}
```

```
input CreateRelationInput {
```

```
  fromConceptId: UUID!
```

```
  toConceptId: UUID!
```

```
  relationType: RelationType!
```

```
  strength: UnitFloat
```

```
  evidence: String
```

```
}
```

```
""""
```

All relationship/edge types in the Apache AGE knowledge graph.

Maps directly to edge labels defined in the graph ontology.

```
""""
```

```
enum RelationType {
```

```
  """Semantic similarity or co-occurrence relationship"""
```

```
  RELATED_TO
```

```
  """Concept A is a prerequisite for understanding Concept B"""
```

```
  PREREQUISITE_OF
```

```
  """Concept B was derived from Concept A"""
```

```
  DERIVED_FROM
```

```
  """Source cites a concept (with optional page/timestamp)"""
```

```
  CITES
```

```
  """Term/Concept belongs to a TopicCluster"""
```

```
  BELONGS_TO
```

```
  """A transcript segment mentions a concept"""
```

```
  MENTIONS
```

```
  """A source or concept refers to another entity"""
```

```
  REFERS_TO
```

```
  """Two concepts contradict each other"""
```

```
  CONTRADICTS
```

```
  """A source was authored by a person"""
```

```
AUTHORED_BY
    """AI-inferred relationship (pending review)"""
INFERRED_RELATED
}
```

```
input CreateContradictionInput {
    conceptIdA: UUID!
    conceptIdB: UUID!
    evidence: String!
    severity: ContradictionSeverity!
    sourceAssetIdA: UUID
    sourceAssetIdB: UUID
}
```

```
input ConceptFilterInput {
    """Full-text search on label, aliases, and description"""
    search: String
    """Filter by domain"""
    domain: String
    """Filter by source type"""
    sourceType: ConceptSourceType
    """Minimum confidence threshold"""
    minConfidence: UnitFloat
}
```

```
input ConceptOrderByInput {
    field: ConceptSortField!
    direction: SortDirection!
}
```

```
enum ConceptSortField {
    LABEL
    CONFIDENCE
    CREATED_AT
    DOMAIN
}
```

---

```
# Connection Types
```

---

```
type ConceptConnection {
    edges: [ConceptEdge!]!
    pageInfo: PageInfo!
    totalCount: NonNegativeInt
}
```

```
type ConceptEdge {
  node: Concept!
  cursor: Cursor!
}

type KnowledgeRelationConnection {
  edges: [KnowledgeRelationEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

type KnowledgeRelationEdge {
  node: KnowledgeRelation!
  cursor: Cursor!
}

type ConceptMentionConnection {
  edges: [ConceptMentionEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

type ConceptMentionEdge {
  node: ConceptMention!
  cursor: Cursor!
}

type SemanticSearchResultConnection {
  edges: [SemanticSearchResultEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}

type SemanticSearchResultEdge {
  node: SemanticSearchResult!
  cursor: Cursor!
}

"""Hybrid search result combining vector similarity with graph context"""
type HybridSearchResult {
  """Vector search result"""
  segment: TranscriptSegment!
  asset: MediaAsset!
  matchedText: String!
  similarity: UnitFloat!
  """Additional context from knowledge graph traversal"""
  graphContext: [ConceptWithRelations!]!
}
```

```
"""Combined relevance score (vector + graph)"""
combinedScore: Float!
```

```
}
```

```
type HybridSearchResultConnection {
  edges: [HybridSearchResultEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}
```

```
type HybridSearchResultEdge {
  node: HybridSearchResult!
  cursor: Cursor!
}
```

```
type TopicClusterConnection {
  edges: [TopicClusterEdge!]!
  pageInfo: PageInfo!
  totalCount: NonNegativeInt
}
```

```
type TopicClusterEdge {
  node: TopicCluster!
  cursor: Cursor!
}
```

---

```
# _____
# Person / Term / Source Input Types
# _____
```

```
input CreatePersonInput {
  name: String!
  role: String
  externalUrl: URL
}
```

```
input UpdatePersonInput {
  name: String
  role: String
  externalUrl: URL
}
```

```
input CreateTermInput {
  label: String!
  domain: String!
  definition: String
"""Link to a parent concept"""

```

```
parentConceptId: UUID  
}
```

```
input UpdateTermInput {  
  label: String  
  domain: String  
  definition: String  
  parentConceptId: UUID  
}
```

```
input CreateSourceInput {  
  title: String!  
  type: SourceType!  
  externalUrl: URL  
  """Link to an existing internal media asset"""  
  internalAssetId: UUID  
  isbn: String  
  doi: String  
}
```

```
input UpdateSourceInput {  
  title: String  
  type: SourceType  
  externalUrl: URL  
  internalAssetId: UUID  
  isbn: String  
  doi: String  
}
```

---

```
# _____  
# Person / Term / Source Connection Types  
# _____
```

```
type PersonConnection {  
  edges: [PersonEdge!]!  
  pageInfo: PageInfo!  
  totalCount: NonNegativeInt  
}
```

```
type PersonEdge {  
  node: Person!  
  cursor: Cursor!  
}
```

```
type TermConnection {  
  edges: [TermEdge!]!  
  pageInfo: PageInfo!
```

```
totalCount: NonNegativeInt  
}
```

```
type TermEdge {  
  node: Term!  
  cursor: Cursor!  
}
```

```
type SourceConnection {  
  edges: [SourceEdge!]!  
  pageInfo: PageInfo!  
  totalCount: NonNegativeInt  
}
```

```
type SourceEdge {  
  node: Source!  
  cursor: Cursor!  
}
```

## 13. Subscriptions Contract

Subscriptions are handled via **Hive Gateway's WebSocket passthrough** mode. The gateway forwards WebSocket connections (using the `[graphql-ws]` protocol) directly to the subgraph that defines the subscription. For HTTP-based clients (e.g., React Native), the gateway also supports **Server-Sent Events (SSE)** as a fallback transport.

Each subgraph that defines subscriptions backs them with **NATS JetStream** events.

### Subscription Definitions

```
graphql
```

```
# ----- Defined in their respective subgraph schemas -----
```

```
#-----  
# Content Subgraph Subscriptions  
#-----
```

```
# In: apps/subgraph-content/src/schema.graphql
```

```
type Subscription {
```

```
    """
```

```
    Fired when the transcription status of a media asset changes.
```

```
    Useful for showing progress in the upload flow.
```

```
    """
```

```
    transcriptionStatusChanged(assetId: UUID!): MediaAsset! @authenticated
```

```
    """
```

```
    Fired when a new transcript segment is available during live transcription.
```

```
    """
```

```
    transcriptSegmentAdded(assetId: UUID!): TranscriptSegment! @authenticated
```

```
}
```

```
#-----  
# Annotation Subgraph Subscriptions  
#-----
```

```
# In: apps/subgraph-annotation/src/schema.graphql
```

```
type Subscription {
```

```
    """
```

```
    Fired when an annotation is created, updated, or deleted on an asset.
```

```
    Used for real-time annotation collaboration.
```

```
    """
```

```
    annotationChanged(
```

```
        assetId: UUID!
```

```
        layers: [AnnotationLayer!]
```

```
    ): AnnotationChangeEvent! @authenticated
```

```
}
```

```
type AnnotationChangeEvent {
```

```
    """Type of change"""
```

```
    changeType: ChangeType!
```

```
    """The annotation that changed (null if deleted)"""
```

```
    annotation: Annotation
```

```
    """ID of the deleted annotation (only for DELETE)"""
```

```
    deletedAnnotationId: UUID
```

```
}
```

```
enum ChangeType {
```

```
CREATED  
UPDATED  
DELETED  
}
```

---

```
# Collaboration Subgraph Subscriptions
```

---

```
# In: apps/subgraph-collaboration/src/schema.graphql
```

```
type Subscription {
```

```
    """
```

Fired when collaborators join/leave or update cursor position.

Used for showing live presence indicators.

```
    """
```

```
collaboratorPresenceChanged(
```

```
    documentId: UUID!
```

```
): PresenceEvent! @authenticated
```

```
}
```

```
type PresenceEvent {
```

```
    eventType: PresenceEventType!
```

```
    session: CollabSession
```

```
    disconnectedUserId: UUID
```

```
}
```

```
enum PresenceEventType {
```

```
    JOINED
```

```
    LEFT
```

```
    CURSOR_MOVED
```

```
}
```

---

```
# Agent Subgraph Subscriptions
```

---

```
# In: apps/subgraph-agent/src/schema.graphql
```

```
type Subscription {
```

```
    """
```

Fired when an agent execution's status changes.

Stream updates from QUEUED → RUNNING → COMPLETED/FAILED.

Also streams partial output during RUNNING state.

```
    """
```

```
agentExecutionUpdated(
```

```
    executionId: UUID!
```

```
): AgentExecution! @authenticated
```

```
"""
Stream the agent's response tokens as they're generated.  
Enables showing the agent's response character by character.
```

```
"""
agentResponseStream(  
    executionId: UUID!  
) : AgentStreamChunk! @authenticated  
}
```

```
"""
A chunk of the agent's streaming response"""

```

```
type AgentStreamChunk {  
    """
    The text chunk"""  
    text: String!  
    """
    Whether this is the final chunk"""  
    done: Boolean!  
    """
    Running token count"""  
    tokensSoFar: NonNegativeInt!  
}
```

```
#-----  
# Knowledge Subgraph Subscriptions  
#-----
```

```
# In: apps/subgraph-knowledge/src/schema.graphql  
type Subscription {  
    """
    Fired when new concepts are extracted from content  
(e.g., after transcription completes and NLP pipeline runs).  
    """  
    conceptsExtracted(  
        assetId: UUID!  
    ) : [Concept!]! @authenticated  
}
```

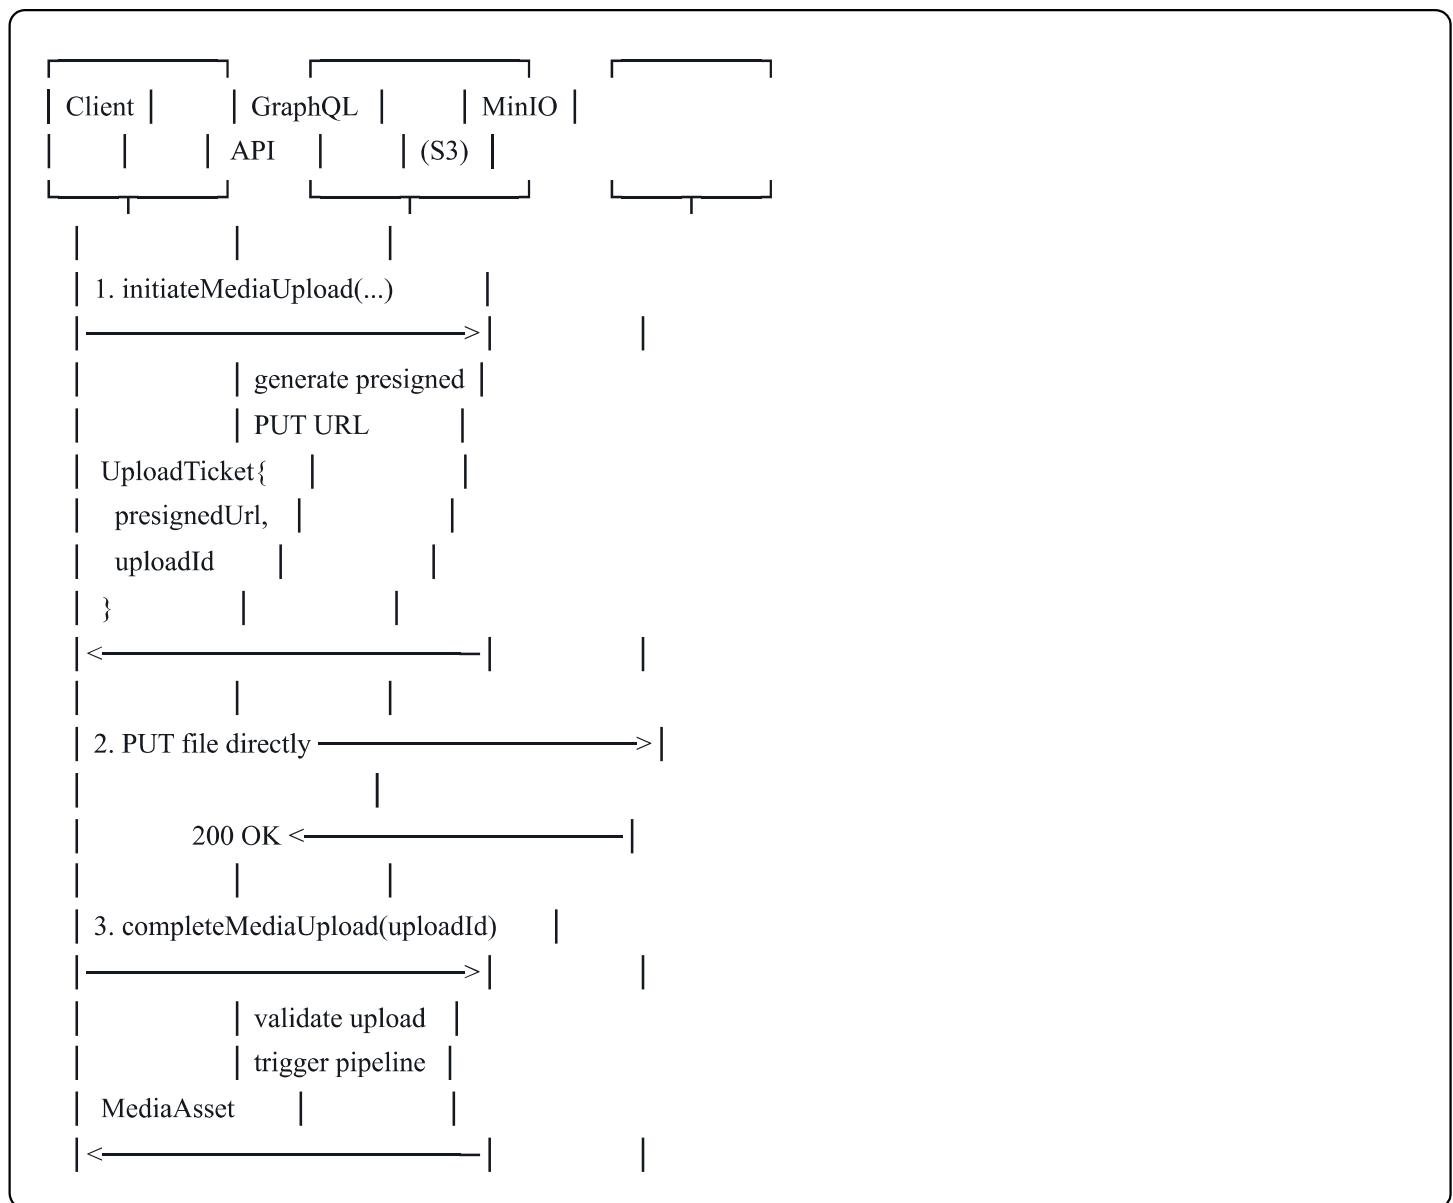
## NATS Subject Mapping

Subscription	NATS Subject Pattern	Publisher
transcriptionStatusChanged	edusphere.{tenant_id}.media.{asset_id}.transcription.status	Media Worker
transcriptSegmentAdded	edusphere.{tenant_id}.media.{asset_id}.segment.added	Transcription Worker
annotationChanged	edusphere.{tenant_id}.annotation.{asset_id}.changed	Annotation Service
collaboratorPresenceChanged	edusphere.{tenant_id}.collab.{document_id}.presence	Hocuspocus Server

Subscription	NATS Subject Pattern	Publisher
agentExecutionUpdated	edusphere.{tenant_id}.agent.execution. {execution_id}.updated	Agent Runner
agentResponseStream	edusphere.{tenant_id}.agent.execution. {execution_id}.stream	Agent Runner
conceptsExtracted	edusphere.{tenant_id}.knowledge. {asset_id}.concepts.extracted	NLP Pipeline

## 14. File Upload Contract (Multipart)

File uploads follow the **two-phase upload pattern**:



**Max file sizes per plan:**

Plan	Max File Size	Max Total Storage
Free	100 MB	1 GB
Starter	500 MB	25 GB
Professional	2 GB	100 GB
Enterprise	10 GB	Unlimited

## 15. Gateway Configuration & Router Settings

### Hive Gateway Configuration

```
typescript
```

```
// —— File: apps/gateway/src/gateway.config.ts ——
```

```
import { defineConfig } from '@graphql-hive/gateway';

export const gatewayConfig = defineConfig({
  // Supergraph SDL source
  supergraph: './supergraph.graphql',

  // Server settings
  port: 4000,
  hostname: '0.0.0.0',
  graphqlEndpoint: '/graphql',
  landingPage: process.env.NODE_ENV !== 'production',

  // CORS
  cors: {
    origin: [
      'https://app.edusphere.dev',
      /\.edusphere\.dev$/,
      'http://localhost:3000',
    ],
    credentials: true,
    allowedHeaders: [
      'Content-Type',
      'Authorization',
      'x-request-id',
      'x-correlation-id',
    ],
  },
}

// JWT Authentication (Keycloak)
jwt: {
  jwks: {
    url: 'https://keycloak.edusphere.dev/realms/edusphere/protocol/openid-connect/certs',
  },
  tokenLookupLocations: [
    { type: 'header', name: 'authorization', prefix: 'Bearer' },
  ],
  // Claims are extracted and propagated as headers to subgraphs
  forward: {
    payload: true, // Forward full JWT payload
    extensions: true,
  },
}

// Subscription support (WebSocket passthrough)
```

```
subscriptions: {  
    enabled: true,  
    protocol: 'ws',      // graphql-ws protocol  
    path: '/ws',  
},  
  
// Header propagation to subgraphs  
propagateHeaders: {  
    fromClientToSubgraphs({ request, subgraphName }) {  
        return {  
            authorization: request.headers.get('authorization') ?? "",  
            'x-tenant-id': request.headers.get('x-tenant-id') ?? "",  
            'x-user-id': request.headers.get('x-user-id') ?? "",  
            'x-user-role': request.headers.get('x-user-role') ?? "",  
            'x-request-id': request.headers.get('x-request-id') ?? crypto.randomUUID(),  
            'x-correlation-id': request.headers.get('x-correlation-id') ?? "",  
            'x-gateway-version': '1.0.0',  
        };  
    },  
},  
  
// OpenTelemetry integration  
openTelemetry: {  
    exporters: [  
        {  
            type: 'otlp',  
            endpoint: 'http://jaeger:4317',  
        },  
    ],  
},  
  
// Prometheus metrics  
prometheus: {  
    enabled: true,  
    endpoint: '/metrics',  
},  
  
// Rate limiting  
rateLimiting: {  
    global: {  
        max: 10000,  
        window: '1s',  
    },  
},  
  
// Query depth limiting  
security: {
```

```
maxDepth: 15,  
maxAliases: 30,  
maxDirectives: 50,  
maxTokens: 10000,  
},  
  
// Persisted operations (APQ)  
persistedOperations: {  
    enabled: true,  
    allowArbitraryOperations: process.env.NODE_ENV !== 'production',  
},  
  
// Execution settings  
executionCancellation: true,  
maskedErrors: process.env.NODE_ENV === 'production',  
includeSubgraphErrors: process.env.NODE_ENV !== 'production',  
});
```

## Subgraph NestJS Bootstrap (GraphQL Yoga)

typescript

```

// — File: apps/subgraph-core/src/app.module.ts —
// Pattern repeated for each subgraph with its own schema

import { Module } from '@nestjs/common';
import { GraphQLModule } from '@graphql-yoga/nestjs';
import { YogaFederationDriver, YogaFederationDriverConfig } from '@graphql-yoga/nestjs/federation';

@Module({
  imports: [
    GraphQLModule.forRoot<YogaFederationDriverConfig>({
      driver: YogaFederationDriver,
      typePaths: ['./src/**/*.graphql'],
      // GraphQL Yoga plugins
      plugins: [],
      // Context extraction from gateway headers
      context: ({ request }) => ({
        tenantId: request.headers.get('x-tenant-id'),
        userId: request.headers.get('x-user-id'),
        userRole: request.headers.get('x-user-role'),
        requestId: request.headers.get('x-request-id'),
      }),
    }),
    // ... domain modules
  ],
})
export class AppModule {}

```

## Supergraph Composition

bash

```
# Compose supergraph SDL from subgraph schemas using Hive CLI
# Install: npm install -g @graphql-hive/cli

# Local composition (dev)
hive supergraph compose \
--config supergraph-config.yaml \
--output supergraph.graphql

# Publish schema to Hive registry (CI/CD)
hive schema:publish \
--registry.endpoint https://app.graphql-hive.com \
--registry.accessToken $HIVE_TOKEN \
--service core \
--url http://subgraph-core:4001/graphql \
apps/subgraph-core/src/schema.graphql

# Schema check (CI — detects breaking changes)
hive schema:check \
--registry.endpoint https://app.graphql-hive.com \
--registry.accessToken $HIVE_TOKEN \
--service core \
apps/subgraph-core/src/schema.graphql
```

yaml

```

# ----- File: supergraph-config.yaml -----
# Hive Gateway supergraph composition config
# Compatible with Apollo Federation v2.7 SDL spec

subgraphs:
  core:
    url: http://subgraph-core:4001/graphql
    schema:
      file: apps/subgraph-core/src/schema.graphql
  content:
    url: http://subgraph-content:4002/graphql
    schema:
      file: apps/subgraph-content/src/schema.graphql
  annotation:
    url: http://subgraph-annotation:4003/graphql
    schema:
      file: apps/subgraph-annotation/src/schema.graphql
  collaboration:
    url: http://subgraph-collaboration:4004/graphql
    schema:
      file: apps/subgraph-collaboration/src/schema.graphql
  agent:
    url: http://subgraph-agent:4005/graphql
    schema:
      file: apps/subgraph-agent/src/schema.graphql
  knowledge:
    url: http://subgraph-knowledge:4006/graphql
    schema:
      file: apps/subgraph-knowledge/src/schema.graphql

```

## 16. Schema Governance & Evolution Rules

### Breaking Change Policy

Change Type	Allowed?	Process
Add field to type	<input checked="" type="checkbox"/> Always	PR review
Add optional argument	<input checked="" type="checkbox"/> Always	PR review
Add new type/enum value	<input checked="" type="checkbox"/> Always	PR review
Add required argument	 With default	PR + migration guide
Rename field	 Never	Add new field, deprecate old

Change Type	Allowed?	Process
Remove field	✗ Never (immediately)	Deprecate → 90-day sunset
Change field type	✗ Never	Add new field, deprecate old
Remove enum value	✗ Never (immediately)	Deprecate → 90-day sunset
Change nullability (nullable → non-null)	✗ Output fields	PR review for input fields
Change nullability (non-null → nullable)	✓ Output fields	PR review

## Deprecation Lifecycle

1. PR submitted with `@deprecated(reason: "Use xyz instead", removalDate: "2026-01-01")`
2. Schema check detects usage in persisted queries → notify consuming teams
3. Field continues to work for 90 days
4. After removalDate: field is removed if usage drops to 0
5. If usage > 0 at removalDate: extend by 30 days, notify teams

## Naming Conventions

Element	Convention	Example
Types	PascalCase	<code>MediaAsset</code> , <code>AgentExecution</code>
Fields	camelCase	<code>createdAt</code> , <code>syllabusJson</code>
Enums	UPPER_SNAKE_CASE	<code>QUIZ_MASTER</code> , <code>AI_GENERATED</code>
Enum types	PascalCase	<code>AnnotationLayer</code> , <code>MediaType</code>
Input types	PascalCase + <code>Input</code> suffix	<code>CreateCourseInput</code>
Filter inputs	PascalCase + <code>FilterInput</code> suffix	<code>CourseFilterInput</code>
Connection types	PascalCase + <code>Connection</code> suffix	<code>CourseConnection</code>
Edge types	PascalCase + <code>Edge</code> suffix	<code>CourseEdge</code>
Mutations	camelCase verb+noun	<code>createCourse</code> , <code>toggleCoursePublished</code>
Queries (single)	camelCase noun	<code>course</code> , <code>user</code>
Queries (list)	camelCase plural noun	<code>courses</code> , <code>users</code>
Subscriptions	camelCase event description	<code>annotationChanged</code> , <code>agentResponseStream</code>

## 17. Client Operation Examples

### Initialize App (Single Request)

```
graphql

query InitializeApp {
  me {
    id
    email
    displayName
    avatarUrl
    role
    preferences {
      language
      theme
      defaultAnnotationLayer
    }
    tenant {
      id
      name
      slug
      plan
      settings {
        defaultLanguage
        allowedLlmProviders
        brandingColors {
          primary
          secondary
        }
      }
    }
  }
}
```

### Browse Course with Video Player Context

```
graphql
```

```
query CoursePlayerView($courseId: UUID!, $assetId: UUID!, $timestamp: Float!) {
  course(id: $courseId) {
    id
    title
    modules(first: 50, orderBy: { field: ORDER_INDEX, direction: ASC }) {
      edges {
        node {
          id
          title
          orderIndex
          mediaAssets(first: 20) {
            edges {
              node {
                id
                title
                type
                durationSeconds
                thumbnailUrl
              }
            }
          }
        }
      }
    }
  }

  mediaAsset(id: $assetId) {
    id
    title
    hlsManifestUrl
    durationSeconds
    transcript {
      id
      fullText
      language
    }
    segments(first: 50) {
      edges {
        node {
          id
          startTime
          endTime
          text
          speaker
        }
      }
    }
  }
}
```

```
}

annotations(
  first: 30
  filter: { layers: [PERSONAL, SHARED, INSTRUCTOR, AI_GENERATED] }
) {
  edges {
    node {
      id
      type
      layer
      content
      timestampStart
      timestampEnd
      author {
        id
        displayName
        avatarUrl
      }
    }
  }
}

semanticContextAt(timestamp: $timestamp) {
  activeConcepts {
    concept {
      id
      label
      domain
    }
    connections {
      relatedConcept {
        id
        label
      }
      relationType
      strength
    }
  }
  activeContradictions {
    concept { id label }
    evidence
    severity
  }
  suggestedPrerequisites {
    id
    label
  }
}
```

```
}
```

## Execute Agent with Streaming

```
graphql

mutation RunChavruta($agentId: UUID!, $annotationId: UUID!) {
  executeAgent(input: {
    agentId: $agentId
    input: { annotationId: $annotationId, mode: "debate" }
  }) {
    id
    status
  }
}

subscription StreamAgentResponse($executionId: UUID!) {
  agentResponseStream(executionId: $executionId) {
    text
    done
    tokensSoFar
  }
}
```

## Semantic Search

```
graphql
```

```
query SmartSearch($query: String!) {
  hybridSearch(query: $query, first: 10, graphDepth: 2) {
    edges {
      node {
        segment {
          id
          text
          startTime
          endTime
        }
        asset {
          id
          title
          type
          thumbnailUrl
        }
        matchedText
        similarity
        combinedScore
      }
      graphContext {
        concept {
          id
          label
          domain
        }
        connections {
          relatedConcept { id label }
          relationType
          strength
        }
      }
    }
  }
  pageInfo {
    hasNextPage
    endCursor
  }
  totalCount
}
}
```

## 18. TypeScript Codegen Configuration

### GraphQL Code Generator Setup

typescript

```
// —— File: packages/graphql-codegen/codegen.ts ——
```

```
import { CodegenConfig } from '@graphql-codegen/cli';

const config: CodegenConfig = {
  schema: [
    'apps/subgraph-core/src/schema.graphql',
    'apps/subgraph-content/src/schema.graphql',
    'apps/subgraph-annotation/src/schema.graphql',
    'apps/subgraph-collaboration/src/schema.graphql',
    'apps/subgraph-agent/src/schema.graphql',
    'apps/subgraph-knowledge/src/schema.graphql',
    'packages/graphql-shared/src/**/*.graphql',
  ],
  documents: [
    'apps/web/src/**/*.graphql',
    'apps/mobile/src/**/*.graphql',
  ],
  generates: {
    // —— Shared types for all packages ——
    'packages/graphql-types/src/generated/types.ts': {
      plugins: [
        'typescript',
        'typescript-operations',
      ],
      config: {
        scalars: {
          UUID: 'string',
          DateTime: 'string',
          JSON: 'Record<string, unknown>',
          JSONObject: 'Record<string, unknown>',
          Base64: 'string',
          Cursor: 'string',
          URL: 'string',
          EmailAddress: 'string',
          NonNegativeInt: 'number',
          PositiveInt: 'number',
          UnitFloat: 'number',
        },
        enumsAsTypes: false,
        avoidOptionals: false,
        maybeValue: 'T | null',
        strictScalars: true,
        useTypeImports: true,
      },
    },
  },
};
```

```
// —— React hooks for web app (using graphql-request or urql) ——  
'apps/web/src/generated/graphql.ts': {  
  preset: 'client',  
  plugins: [  
    'typescript',  
    'typescript-operations',  
    'typescript-react-query', // TanStack Query hooks  
  ],  
  config: {  
    fetcher: {  
      endpoint: 'process.env.VITE_GRAPHQL_ENDPOINT',  
      fetchParams: {  
        headers: {  
          'content-type': 'application/json',  
        },  
      },  
    },  
    exposeFetcher: true,  
    exposeQueryKeys: true,  
    addSuspenseQuery: true,  
  },  
},
```

// —— NestJS resolver types for each subgraph ——

```
'apps/subgraph-core/src/generated/resolvers.ts': {  
  plugins: ['typescript', 'typescript-resolvers'],  
  config: {  
    federation: true,  
    mappers: {  
      Tenant: '../entities/tenant.entity#TenantEntity',  
      User: '../entities/user.entity#UserEntity',  
    },  
    contextType: '../context#GraphQLContext',  
  },  
},  
'apps/subgraph-content/src/generated/resolvers.ts': {  
  plugins: ['typescript', 'typescript-resolvers'],  
  config: {  
    federation: true,  
    mappers: {  
      Course: '../entities/course.entity#CourseEntity',  
      Module: '../entities/module.entity#ModuleEntity',  
      MediaAsset: '../entities/media-asset.entity#MediaAssetEntity',  
      Transcript: '../entities/transcript.entity#TranscriptEntity',  
      TranscriptSegment: '../entities/segment.entity#SegmentEntity',  
    },  
  },
```

```

    contextType: '../context#GraphQLContext',
  },
},
// ... (similar for annotation, collaboration, agent, knowledge)

// ----- Schema introspection for tooling -----
'packages/graphql-types/src/generated/schema.json': {
  plugins: ['introspection'],
},

// ----- Fragment matcher for client cache -----
'apps/web/src/generated/fragment-matcher.ts': {
  plugins: ['fragment-matcher'],
},
},
hooks: {
  afterAllFileWrite: ['prettier --write'],
},
};

export default config;

```

## NPM Scripts

```

bash

# Generate all types
pnpm --filter @edusphere/graphql-codegen exec graphql-codegen

# Watch mode during development
pnpm --filter @edusphere/graphql-codegen exec graphql-codegen --watch

# Check schema compatibility (via Hive CLI)
hive schema:check \
--registry.endpoint https://app.graphql-hive.com \
--registry.accessToken $HIVE_TOKEN \
--service content \
apps/subgraph-content/src/schema.graphql

```

# Full Query/Mutation/Subscription Summary

## Queries (44 total)

Subgraph	Query	Auth
Core	me	✓
Core	user(id)	✓
Core	users(filter, orderBy, pagination)	✓
Core	currentTenant	✓
Core	tenantBySlug(slug)	Public
Content	course(id)	✓
Content	courses(filter, orderBy, pagination)	✓
Content	module(id)	✓
Content	mediaAsset(id)	✓
Content	mediaAssets(filter, orderBy, pagination)	✓
Content	segmentsForTimeRange(assetId, start, end)	✓
Content	searchTranscripts(query, pagination)	✓
Annotation	annotation(id)	✓
Annotation	annotations(filter, orderBy, pagination)	✓
Annotation	annotationThread(rootId)	✓
Collaboration	collabDocument(id)	✓
Collaboration	collabDocumentByName(name)	✓
Collaboration	collabDocumentsForEntity(type, id)	✓
Collaboration	collabConnectionInfo(docId)	✓
Agent	agentDefinition(id)	✓
Agent	agentDefinitions(filter, orderBy, pagination)	✓
Agent	agentTemplates	✓

Subgraph	Query	Auth
Agent	agentExecution(id)	✓
Agent	agentExecutions(filter, pagination)	✓
Knowledge	concept(id)	✓
Knowledge	concepts(filter, orderBy, pagination)	✓
Knowledge	semanticSearch(query, pagination)	✓
Knowledge	hybridSearch(query, graphDepth)	✓
Knowledge	relatedConcepts(conceptId, maxDepth)	✓
Knowledge	contradictions(conceptId)	✓
Knowledge	learningPath(conceptId, maxDepth)	✓
Knowledge	topicClusters(pagination)	✓
Knowledge	person(id)	✓
Knowledge	people(search, pagination)	✓
Knowledge	term(id)	✓
Knowledge	terms(domain, search, pagination)	✓
Knowledge	source(id)	✓
Knowledge	sources(type, search, pagination)	✓

## Mutations (44 total)

Subgraph	Mutation	Required Scope
Core	updateMyProfile	authenticated
Core	updateUserRole	org:users
Core	deactivateUser	org:users
Core	reactivateUser	org:users
Core	updateTenantSettings	org:manage
Content	createCourse	course:write

Subgraph	Mutation	Required Scope
Content	updateCourse	course:write
Content	deleteCourse	course:write
Content	toggleCoursePublished	course:write
Content	forkCourse	course:write
Content	createModule	course:write
Content	updateModule	course:write
Content	deleteModule	course:write
Content	reorderModules	course:write
Content	initiateMediaUpload	media:upload
Content	completeMediaUpload	media:upload
Content	updateMediaAsset	course:write
Content	deleteMediaAsset	course:write
Content	retriggerTranscription	course:write
Annotation	createAnnotation	authenticated
Annotation	updateAnnotation	authenticated (owner)
Annotation	deleteAnnotation	authenticated (owner)
Annotation	toggleAnnotationPin	annotation:write
Annotation	resolveAnnotation	authenticated
Annotation	moveAnnotationsToLayer	annotation:write
Collaboration	createCollabDocument	authenticated
Collaboration	compactCollabDocument	org:manage
Agent	createAgentDefinition	agent:write
Agent	updateAgentDefinition	agent:write (owner)
Agent	deleteAgentDefinition	agent:write (owner)

Subgraph	Mutation	Required Scope
Agent	executeAgent	agent:execute
Agent	cancelAgentExecution	authenticated
Knowledge	createConcept	knowledge:write
Knowledge	updateConcept	knowledge:write
Knowledge	deleteConcept	knowledge:write
Knowledge	createRelation	knowledge:write
Knowledge	deleteRelation	knowledge:write
Knowledge	createContradiction	knowledge:write
Knowledge	reindexAssetEmbeddings	knowledge:write
Knowledge	reviewInferredRelation	knowledge:write
Knowledge	createPerson	knowledge:write
Knowledge	updatePerson	knowledge:write
Knowledge	deletePerson	knowledge:write
Knowledge	createTerm	knowledge:write
Knowledge	updateTerm	knowledge:write
Knowledge	deleteTerm	knowledge:write
Knowledge	createSource	knowledge:write
Knowledge	updateSource	knowledge:write
Knowledge	deleteSource	knowledge:write
Knowledge	linkAuthorToSource	knowledge:write
Knowledge	linkSourceToConcept	knowledge:write

## Subscriptions (7 total)

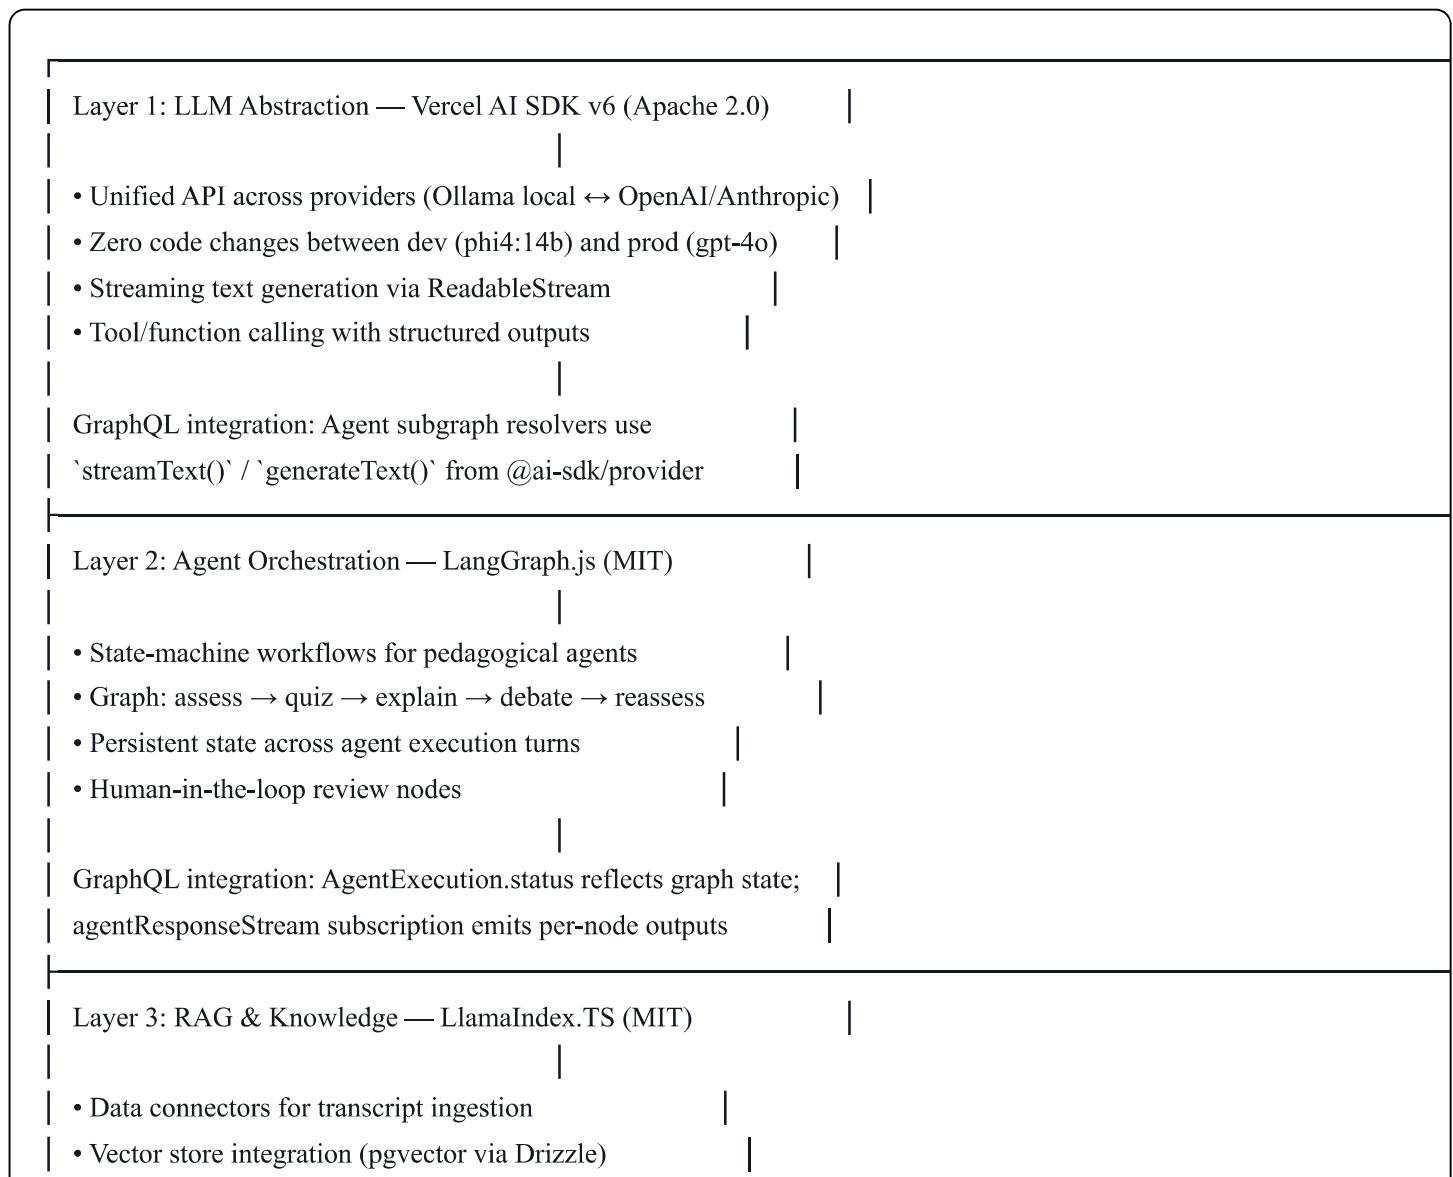
Subgraph	Subscription	Trigger
Content	transcriptionStatusChanged(assetId)	Media pipeline

Subgraph	Subscription	Trigger
Content	transcriptSegmentAdded(assetId)	Transcription worker
Annotation	annotationChanged(assetId, layers)	Annotation CRUD
Collaboration	collaboratorPresenceChanged(documentId)	Hocuspocus
Agent	agentExecutionUpdated(executionId)	Agent runner
Agent	agentResponseStream(executionId)	Agent LLM stream
Knowledge	conceptsExtracted(assetId)	NLP pipeline

## 19. AI/ML Architecture Integration

This section documents how the GraphQL API integrates with the platform's three-layer AI architecture, as specified in the EduSphere Architecture Guide.

### Three-Layer AI Architecture



- Knowledge graph integration (Apache AGE Cypher queries)
- HybridRAG: parallel vector + graph retrieval → fused context

GraphQL integration: hybridSearch query, semanticSearch query, semanticContextAt on MediaAsset

## Agent Template → LangGraph Workflow Mapping

Each `AgentTemplate` enum value maps to a pre-built LangGraph.js state machine:

<b>Template</b>	<b>LangGraph</b>	<b>Description</b>
	<b>Workflow</b>	
<code>CHAVRUTA</code>	<code>chavruta-debate-graph</code>	Dialectical debate: present thesis → find contradictions → argue both sides → synthesize. Uses <code>CONTRADICTS</code> edges from knowledge graph
<code>SUMMARIZER</code>	<code>summarize-graph</code>	Progressive summarization: chunk → summarize → merge → condense. Operates on transcript segments
<code>QUIZ_MASTER</code>	<code>quiz-assess-graph</code>	Adaptive quizzing: assess level → generate questions → evaluate answers → adjust difficulty. Uses <code>PREREQUISITE_OF</code> edges
<code>RESEARCH_SCOUT</code>	<code>research-scout-graph</code>	Cross-reference finder: semantic search → graph traversal → contradiction detection → report. Powers deep research
<code>EXPLAINER</code>	<code>explain-graph</code>	Adaptive explanation: detect knowledge gaps → explain with analogies → verify understanding. Uses prerequisite chains
<code>CUSTOM</code>	User-defined via JSON config	Custom LangGraph state machine loaded from <code>AgentConfig.systemPrompt</code> with configurable nodes

## MCP (Model Context Protocol) Integration

Agents use MCP for tool integrations. The `AgentConfig.toolsEnabled` field specifies which MCP tools an agent can call:

typescript

```
// Available MCP tool categories
type MCPToolCategory =
| 'knowledge_graph' // Query/create concepts, relations, contradictions
| 'semantic_search' // Vector search across content
| 'transcript_reader' // Read transcript segments by time range
| 'annotation_writer' // Create annotations on behalf of user
| 'web_search' // External web search (sandboxed)
| 'calculator' // Mathematical computation
| 'citation_formatter'; // Format academic citations
```

## LLM Provider Routing

The Vercel AI SDK provider system is configured per-tenant:

```
typescript

// Tenant settings control LLM routing
// Maps to TenantSettings.allowedLlmProviders in the database schema

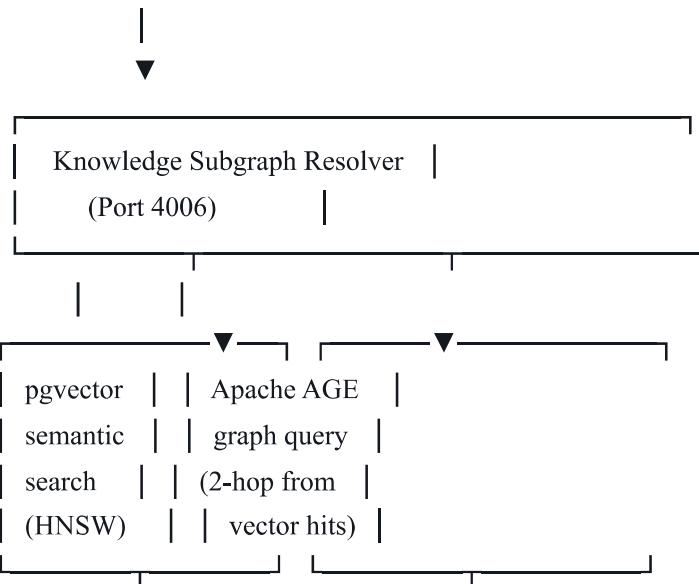
// Development: Ollama (local)
// - Llama 3.1 8B (4-6 GB VRAM) — general tutoring
// - Phi-4 14B (8-10 GB VRAM) — reasoning-heavy explanations

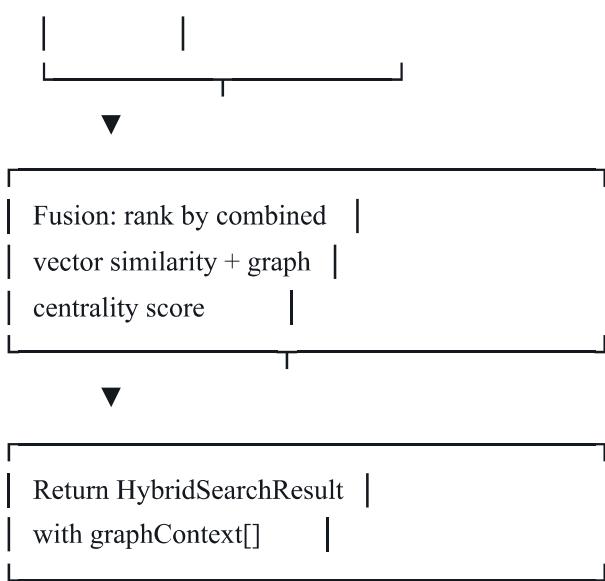
// Production: Cloud providers
// - OpenAI (gpt-4o, gpt-4o-mini)
// - Anthropic (claude-sonnet-4-20250514)

// Agent-level override via AgentConfig.modelOverride
// e.g., "phi4:14b" for a specific agent definition
```

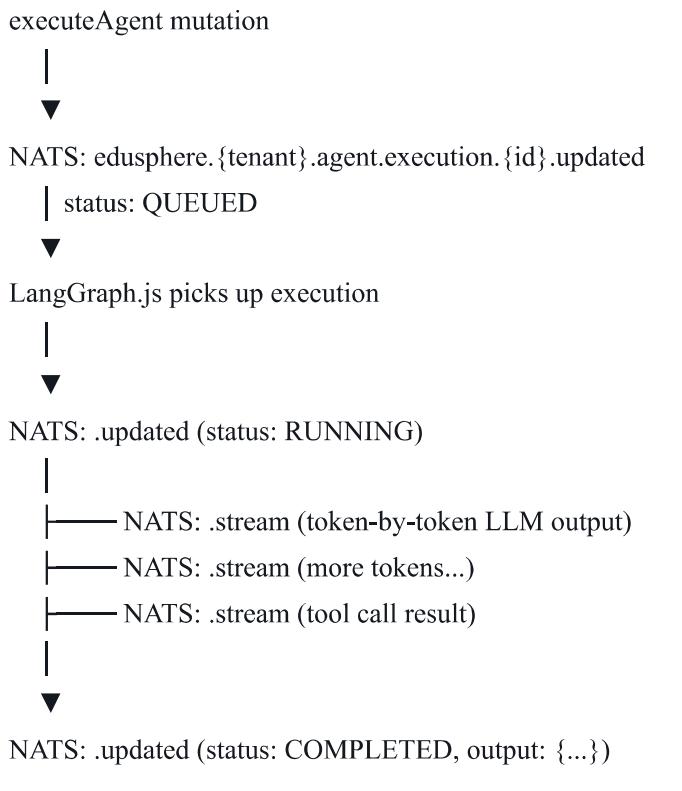
## HybridRAG Search Flow (GraphQL → Database)

Client sends: hybridSearch(query: "Rambam's view on divine attributes")





## Agent Execution Lifecycle (Subscription Flow)



Client subscribes to:

- agentExecutionUpdated(executionId) → status changes
- agentResponseStream(executionId) → real-time token stream

## Agent Sandboxing

For multi-tenant safety, agent executions are sandboxed:

- **gVisor** (Apache 2.0, by Google) provides user-space kernel isolation with only 10-20% performance overhead
- Each agent execution runs in an isolated gVisor container
- Resource limits: CPU, memory, network, and time are capped per-tenant plan

- MCP tool calls are proxy-mediated — agents cannot directly access databases or external services
- 

## 20. Technology Stack Cross-Reference

This table maps every technology choice to its usage in the GraphQL API layer:

Technology	License	GraphQL API Usage
<b>Hive Gateway v2</b>	MIT	Supergraph gateway, query planning, auth enforcement, subscriptions
<b>GraphQL Yoga</b>	MIT	Subgraph HTTP server within each NestJS service
<b>NestJS</b>	MIT	Service framework for all 6 subgraphs
<b>Drizzle ORM</b>	Apache 2.0	Database access in resolvers, RLS via <code>withTenantContext()</code>
<b>PostgreSQL 16+</b>	PostgreSQL	Relational data store with RLS policies
<b>Apache AGE</b>	Apache 2.0	Knowledge graph queries (Cypher via <code>cypher()</code> function)
<b>pgvector</b>	PostgreSQL	Embedding storage and HNSW semantic search
<b>Keycloak v26</b>	Apache 2.0	JWT issuer, JWKS endpoint for gateway validation
<b>NATS JetStream</b>	Apache 2.0	Event transport for subscriptions, async agent execution
<b>Yjs + Hocuspocus</b>	MIT	CRDT collaboration (CollabDocument, WebSocket URLs)
<b>Vercel AI SDK v6</b>	Apache 2.0	LLM calls in Agent subgraph resolvers
<b>LangGraph.js</b>	MIT	Agent workflow orchestration (state machines)
<b>LlamaIndex.TS</b>	MIT	RAG pipeline, knowledge graph indexing
<b>MinIO</b>	AGPLv3	S3-compatible storage for media uploads (presigned URLs)
<b>faster-whisper</b>	MIT	Transcription pipeline (triggers via NATS events)
<b>Video.js v8</b>	Apache 2.0	Client-side: video player consuming HLS manifests
<b>Konva.js v10</b>	MIT	Client-side: sketch annotation canvas rendering
<b>React + Vite</b>	MIT	Web client consuming the supergraph
<b>TanStack Query v5</b>	MIT	Client-side: GraphQL data fetching and caching
<b>Expo SDK 54</b>	MIT	Mobile client with offline-first patterns
<b>GraphQL Hive</b>	MIT	Schema registry, breaking change detection
<b>Traefik v3.6</b>	MIT	Reverse proxy / ingress controller in production

## Database Schema → GraphQL Type Mapping

Database Table	GraphQL Type	Subgraph
tenants	Tenant	Core
users	User	Core
courses	Course	Content
modules	Module	Content
media_assets	MediaAsset	Content
transcripts	Transcript	Content
transcript_segments	TranscriptSegment	Content
annotations	Annotation	Annotation
collab_documents	CollabDocument	Collaboration
crdt_updates	(internal, not exposed)	Collaboration
collab_sessions	CollabSession	Collaboration
agent_definitions	AgentDefinition	Agent
agent_executions	AgentExecution	Agent
content_embeddings	(internal, powers semanticSearch)	Knowledge
annotation_embeddings	(internal, powers semanticSearch)	Knowledge
concept_embeddings	(internal, powers hybridSearch)	Knowledge

## Apache AGE Graph → GraphQL Type Mapping

AGE Vertex/Edge	GraphQL Type	Notes
Concept vertex	Concept	Primary knowledge graph entity
Person vertex	Person	Authors, speakers, historical figures
Term vertex	Term	Domain-specific terms
Source vertex	Source	External references (books, papers, URLs)

AGE Vertex/Edge	GraphQL Type	Notes
[TopicCluster] vertex	TopicCluster	Auto-generated concept groupings
[RELATED_TO] edge	KnowledgeRelation	Returned via [relatedConcepts] query
[CONTRADICTS] edge	Contradiction	Returned via [contradictions] query
[PREREQUISITE_OF] edge	PrerequisiteLink	Returned via [prerequisites] / [dependents]
[MENTIONS] edge	ConceptMention	Segment-to-concept temporal links
[CITES] edge	(via [citedConcepts])	Source-to-concept citations
[AUTHORED_BY] edge	(via [authors])	Person-to-source authorship
[INFERRRED RELATED] edge	KnowledgeRelation	AI-inferred, pending review
[REFERS_TO] edge	KnowledgeRelation	General reference between entities
[DERIVED_FROM] edge	KnowledgeRelation	Concept derivation chain
[BELONGS_TO] edge	(via [topicCluster])	Concept-to-cluster membership