

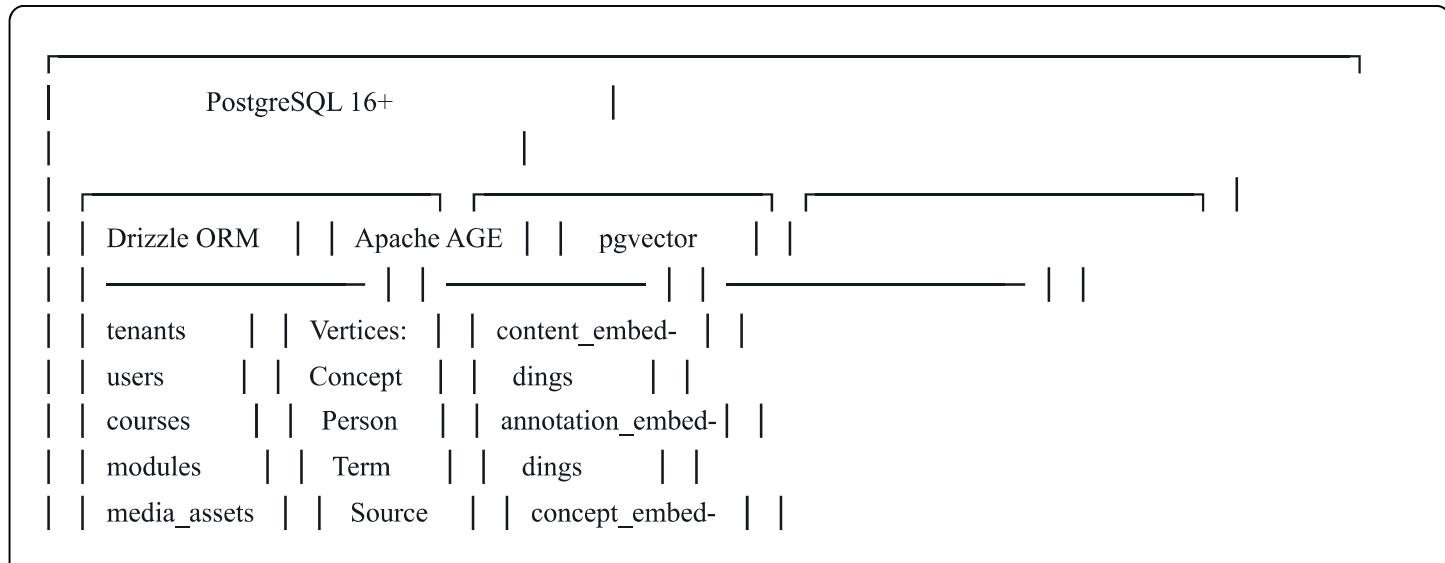
EduSphere Database Schema Definition

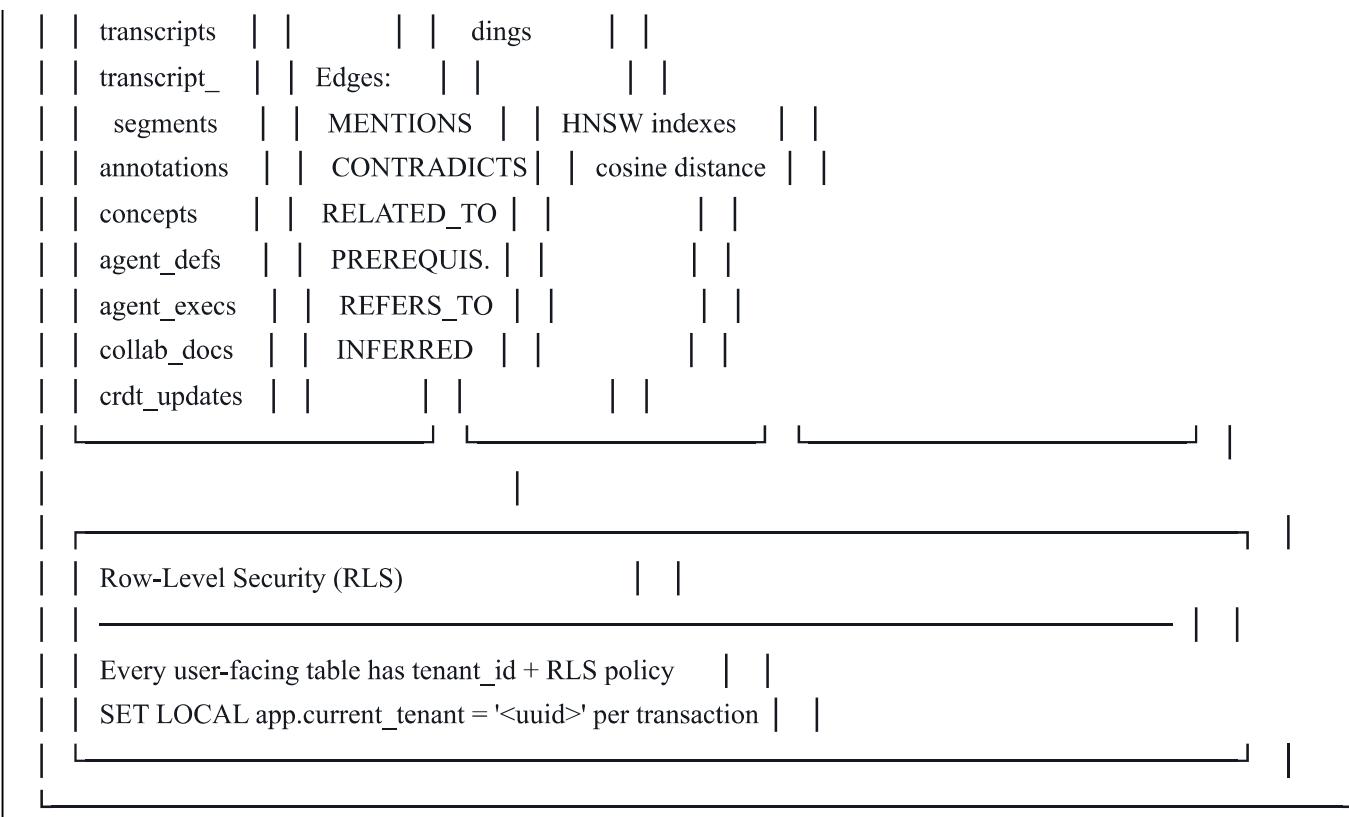
Reference document for Claude Code and all developers. Every table, graph node, edge type, vector index, and RLS policy is defined here. Do NOT deviate from these definitions without updating this document first.

Table of Contents

1. [Architecture Overview](#)
 2. [PostgreSQL Extensions Setup](#)
 3. [Shared Definitions \(Enums, Roles, Helpers\)](#)
 4. [Core Schema — Tenants & Users](#)
 5. [Content Schema — Courses, Media, Transcripts](#)
 6. [Annotation Schema — Layers, Sketches, Comments](#)
 7. [Collaboration Schema — CRDT Persistence](#)
 8. [Agent Schema — Definitions, Executions, Templates](#)
 9. [Row-Level Security Policies](#)
 10. [Apache AGE — Knowledge Graph Ontology](#)
 11. [pgvector — Embedding Tables & Indexes](#)
 12. [Drizzle Relations \(RQBv2\)](#)
 13. [Migration Strategy](#)
 14. [Seed Data](#)
 15. [TypeScript Graph Query Helpers](#)
-

1. Architecture Overview





Key Design Decisions:

- `uuid` for all primary keys (globally unique, no tenant collision)
- `tenant_id` on every user-facing table (RLS enforced)
- Timestamps use `timestamp with time zone` (always UTC)
- Soft deletes via `deleted_at` column (never hard delete user data)
- `jsonb` for flexible/schemaless fields (syllabus, config, metadata)
- Apache AGE graph name: `edusphere_graph` (one graph per database)
- pgvector dimensions: 768 (nomic-embed-text) — configurable via constant

2. PostgreSQL Extensions Setup

sql

```

-- File: packages/db/src/migrations/0000_extensions.sql
-- Run ONCE as superuser before Drizzle migrations

-- Required extensions
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";          -- UUID generation
CREATE EXTENSION IF NOT EXISTS "pgcrypto";            -- Crypto functions
CREATE EXTENSION IF NOT EXISTS "age";                 -- Apache AGE graph
CREATE EXTENSION IF NOT EXISTS "vector";              -- pgvector

-- Load AGE into search path (required per-connection)
-- This is handled in the connection pool initialization
LOAD 'age';
SET search_path = ag_catalog, "$user", public;

-- Create the knowledge graph
SELECT create_graph('edusphere_graph');

-- Create application role for RLS
DO $$

BEGIN
    IF NOT EXISTS (SELECT FROM pg_roles WHERE rolname = 'edusphere_app') THEN
        CREATE ROLE edusphere_app NOLOGIN;
    END IF;
END
$$;

-- Grant necessary permissions
GRANT USAGE ON SCHEMA ag_catalog TO edusphere_app;
GRANT SELECT ON ALL TABLES IN SCHEMA ag_catalog TO edusphere_app;

```

Custom PostgreSQL Docker Image ([\(infrastructure/docker/postgres-age/Dockerfile\)](#)):

dockerfile

```
FROM postgres:16-bookworm

# Install build dependencies
RUN apt-get update && apt-get install -y \
    build-essential libreadline-dev zlib1g-dev flex bison \
    git postgresql-server-dev-16 \
    && rm -rf /var/lib/apt/lists/*

# Install Apache AGE
RUN git clone --branch release/PG16/1.5.0 https://github.com/apache/age.git /tmp/age \
    && cd /tmp/age && make install && rm -rf /tmp/age

# Install pgvector
RUN git clone --branch v0.8.0 https://github.com/pgvector/pgvector.git /tmp/pgvector \
    && cd /tmp/pgvector && make install && rm -rf /tmp/pgvector

# Configure shared_preload_libraries
RUN echo "shared_preload_libraries = 'age'" >> /usr/share/postgresql/postgresql.conf.sample
```

3. Shared Definitions

typescript

```
// File: packages/db/src/schema/_shared.ts

import { sql } from 'drizzle-orm';
import {
  uuid,
  timestamp,
  pgRole,
  pgEnum,
} from 'drizzle-orm/pg-core';

// -----
// Application Role (for RLS policies)
// -----
export const appRole = pgRole('edusphere_app').existing();

// -----
// Enums
// -----


export const userRoleValues = ['super_admin', 'org_admin', 'instructor', 'student', 'researcher'] as const;
export type UserRole = typeof userRoleValues[number];
export const userRoleEnum = pgEnum('user_role', userRoleValues);

export const mediaTypeValues = ['video', 'audio', 'pdf', 'image', 'document'] as const;
export type MediaType = typeof mediaTypeValues[number];
export const mediaTypeEnum = pgEnum('media_type', mediaTypeValues);

export const transcriptionStatusValues = ['pending', 'processing', 'completed', 'failed'] as const;
export type TranscriptionStatus = typeof transcriptionStatusValues[number];
export const transcriptionStatusEnum = pgEnum('transcription_status', transcriptionStatusValues);

export const annotationTypeValues = ['text', 'sketch', 'link', 'bookmark', 'spatial_comment'] as const;
export type AnnotationType = typeof annotationTypeValues[number];
export const annotationTypeEnum = pgEnum('annotation_type', annotationTypeValues);

export const annotationLayerValues = ['personal', 'shared', 'instructor', 'ai_generated'] as const;
export type AnnotationLayer = typeof annotationLayerValues[number];
export const annotationLayerEnum = pgEnum('annotation_layer', annotationLayerValues);

export const agentTemplateValues = ['chavruta', 'summarizer', 'quiz_master', 'research_scout', 'explainer', 'custom'] as const;
export type AgentTemplate = typeof agentTemplateValues[number];
export const agentTemplateEnum = pgEnum('agent_template', agentTemplateValues);

export const agentTriggerValues = ['manual', 'on_annotation', 'on_new_content', 'scheduled'] as const;
export type AgentTrigger = typeof agentTriggerValues[number];
export const agentTriggerEnum = pgEnum('agent_trigger', agentTriggerValues);
```

```

export const agentOutputFormatValues = ['chat', 'overlay', 'side_panel', 'notification'] as const;
export type AgentOutputFormat = typeof agentOutputFormatValues[number];
export const agentOutputFormatEnum = pgEnum('agent_output_format', agentOutputFormatValues);

export const executionStatusValues = ['queued', 'running', 'completed', 'failed', 'cancelled'] as const;
export type ExecutionStatus = typeof executionStatusValues[number];
export const executionStatusEnum = pgEnum('execution_status', executionStatusValues);

export const tenantPlanValues = ['free', 'starter', 'professional', 'enterprise'] as const;
export type TenantPlan = typeof tenantPlanValues[number];
export const tenantPlanEnum = pgEnum('tenant_plan', tenantPlanValues);

// -----
// Shared Column Helpers
// -----


/** Standard primary key: UUID v4, auto-generated */
export const pk = () => uuid('id').primaryKey().defaultRandom();

/** Tenant isolation column — REQUIRED on all user-facing tables */
export const tenantId = () => uuid('tenant_id').notNull();

/** Standard timestamps */
export const timestamps = () => ({
  createdAt: timestamp('created_at', { withTimezone: true }),
  .notNull(),
  .defaultNow(),
  updatedAt: timestamp('updated_at', { withTimezone: true }),
  .notNull(),
  .defaultNow(),
  .$onUpdate(() => new Date()),
});

/** Soft delete column */
export const softDelete = () => ({
  deletedAt: timestamp('deleted_at', { withTimezone: true }),
});

// -----
// RLS Helper SQL fragments
// -----


/** RLS USING clause: tenant_id matches current_setting */
export const rlsTenantCheck = sql`tenant_id = current_setting('app.current_tenant', true)::uuid`;

/** RLS WITH CHECK clause: same as using for inserts */

```

```
export const rlsTenantInsertCheck = sql`tenant_id = current_setting('app.current_tenant', true)::uuid`;  
  
// _____  
// Constants  
// _____  
  
/** pgvector embedding dimensions (nomic-embed-text = 768) */  
export const EMBEDDING_DIMENSIONS = 768;  
  
/** Apache AGE graph name */  
export const GRAPH_NAME = 'edusphere_graph';
```

4. Core Schema — Tenants & Users

typescript

```
// File: packages/db/src/schema/tenants.ts

import { pgTable, pgPolicy, varchar, jsonb, boolean, integer } from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import { pk, tenantId, timestamps, softDelete, tenantPlanEnum, appRole, rlsTenantCheck, rlsTenantInsertCheck } from './shared';

// -----
// tenants — Top-level organizational units
// -----
// NOTE: tenants table does NOT have tenant_id (it IS the tenant)
// RLS is based on the tenant's own id
export const tenants = pgTable('tenants', {
  id: pk(),
  name: varchar('name', { length: 255 }).notNull(),
  slug: varchar('slug', { length: 100 }).notNull().unique(),
  plan: tenantPlanEnum('plan').notNull().default('free'),
  isActive: boolean('is_active').notNull().default(true),
  maxUsers: integer('max_users').notNull().default(50),
  settings: jsonb('settings').$type<TenantSettings>().default({}),
  keycloakOrgId: varchar('keycloak_org_id', { length: 255 }).unique(),
  ...timestamps(),
  ...softDelete(),
}, (table) => [
  pgPolicy('tenants_select_own', {
    for: 'select',
    to: appRole,
    using: sql`(${table.id} = current_setting('app.current_tenant', true)::uuid)`,
  }),
  pgPolicy('tenants_update_own', {
    for: 'update',
    to: appRole,
    using: sql`(${table.id} = current_setting('app.current_tenant', true)::uuid)`,
    withCheck: sql`(${table.id} = current_setting('app.current_tenant', true)::uuid)`,
  }),
  // INSERT and DELETE managed by super_admin only (no RLS policy = denied)
]);
};

/** Tenant settings JSON structure */
export interface TenantSettings {
  defaultLanguage?: string;
  allowedLlmProviders?: string[];
  maxStorageGb?: number;
  customOntologyFields?: Array<{ name: string; type: string }>;
  brandingColors?: { primary: string; secondary: string };
}
```

typescript

```
// File: packages/db/src/schema/users.ts

import { pgTable, pgPolicy, varchar, jsonb, boolean, index } from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {

pk, tenantId, timestamps, softDelete,
userRoleEnum, appRole, rlsTenantCheck, rlsTenantInsertCheck,
} from './_shared.ts';
import { tenants } from './tenants.ts';

// -----
// users — Authenticated users (synced from Keycloak)
// -----


export const users = pgTable('users', {
id: pk(),
tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
keycloakUserId: varchar('keycloak_user_id', { length: 255 }).notNull().unique(),
email: varchar('email', { length: 320 }).notNull(),
displayName: varchar('display_name', { length: 255 }).notNull(),
avatarUrl: varchar('avatar_url', { length: 2048 }),
role: userRoleEnum('role').notNull().default('student'),
isActive: boolean('is_active').notNull().default(true),
preferences: jsonb('preferences').$type<UserPreferences>().default({}),
offlineSyncToken: varchar('offline_sync_token', { length: 500 }),
...timestamps(),
...softDelete(),
}, (table) => [
// Indexes
index('idx_users_tenant').on(table.tenantId),
index('idx_users_email').on(table.email),
index('idx_users_keycloak').on(table.keycloakUserId),

// RLS Policies
pgPolicy('users_select_same_tenant', {
for: 'select',
to: appRole,
using: rlsTenantCheck,
}),
pgPolicy('users_insert_same_tenant', {
for: 'insert',
to: appRole,
withCheck: rlsTenantInsertCheck,
}),
pgPolicy('users_update_same_tenant', {
for: 'update',
to: appRole,
})]
```

```
using: rlsTenantCheck,  
withCheck: rlsTenantInsertCheck,  
},  
];  
  
export interface UserPreferences {  
language?: string;  
theme?: 'light' | 'dark' | 'system';  
defaultAnnotationLayer?: string;  
notificationsEnabled?: boolean;  
agentDefaults?: { difficulty?: string; personality?: string };  
}
```

5. Content Schema — Courses, Media, Transcripts

typescript

```

// File: packages/db/src/schema/courses.ts

import { pgTable, pgPolicy, varchar, text, jsonb, integer, boolean, index } from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import { pk, tenantId, timestamps, softDelete, appRole, rlsTenantCheck, rlsTenantInsertCheck } from './_shared.ts';
import { tenants } from './tenants.ts';
import { users } from './users.ts';

// -----
// courses
// -----


export const courses = pgTable('courses', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  creatorId: pk().references(() => users.id, { onDelete: 'set null' }),
  title: varchar('title', { length: 500 }).notNull(),
  description: text('description'),
  syllabusJson: jsonb('syllabus_json').$type<SyllabusItem[]>().default([]),
  version: integer('version').notNull().default(1),
  isPublished: boolean('is_published').notNull().default(false),
  forkedFromId: pk().references((): any => courses.id, { onDelete: 'set null' }),
  tags: jsonb('tags').$type<string[]>().default([]),
  ...timestamps(),
  ...softDelete(),
}, (table) => [
  index('idx_courses_tenant').on(table.tenantId),
  index('idx_courses_creator').on(table.creatorId),

  pgPolicy('courses_select_same_tenant', {
    for: 'select', to: appRole, using: rlsTenantCheck,
  }),
  pgPolicy('courses_insert_same_tenant', {
    for: 'insert', to: appRole, withCheck: rlsTenantInsertCheck,
  }),
  pgPolicy('courses_update_same_tenant', {
    for: 'update', to: appRole,
    using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
  }),
  pgPolicy('courses_delete_same_tenant', {
    for: 'delete', to: appRole, using: rlsTenantCheck,
  }),
];

export interface SyllabusItem {
  title: string;
  description?: string;
}

```

```
moduleIds?: string[];
order: number;
}

// -----
// modules — Subdivisions within a course
// -----
```

```
export const modules = pgTable('modules', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  courseId: pk().references(() => courses.id, { onDelete: 'cascade' }).notNull(),
  title: varchar('title', { length: 500 }).notNull(),
  description: text('description'),
  orderIndex: integer('order_index').notNull().default(0),
  ...timestamps(),
}, (table) => [
  index('idx_modules_tenant').on(table.tenantId),
  index('idx_modules_course').on(table.courseId),

  pgPolicy('modules_tenant_isolation', {
    for: 'all', to: appRole,
    using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
  }),
]);
```

typescript

```

// File: packages/db/src/schema/media-assets.ts

import {
  pgTable, pgPolicy, varchar, text, jsonb,
  integer, real, index, uniqueIndex,
} from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {
  pk, tenantId, timestamps, softDelete,
  mediaTypeEnum, transcriptionStatusEnum,
  appRole, rlsTenantCheck, rlsTenantInsertCheck,
} from './_shared.ts';
import { tenants } from './tenants.ts';
import { users } from './users.ts';
import { modules } from './courses.ts';

// -----
// media_assets — Videos, PDFs, audio files, etc.
// -----


export const mediaAssets = pgTable('media_assets', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  uploaderId: pk().references(() => users.id, { onDelete: 'set null' }),
  moduleId: pk().references(() => modules.id, { onDelete: 'set null' }),
  type: mediaTypeEnum('type').notNull(),
  title: varchar('title', { length: 500 }).notNull(),
  description: text('description'),
  s3Path: varchar('s3_path', { length: 2048 }).notNull(),
  s3Bucket: varchar('s3_bucket', { length: 255 }).notNull(),
  originalFilename: varchar('original_filename', { length: 500 }),
  mimeType: varchar('mime_type', { length: 255 }),
  fileSizeBytes: integer('file_size_bytes'),
  durationSeconds: real('duration_seconds'),
  transcriptionStatus: transcriptionStatusEnum('transcription_status').default('pending'),
  hlsManifestPath: varchar('hls_manifest_path', { length: 2048 }),
  thumbnailPath: varchar('thumbnail_path', { length: 2048 }),
  metadata: jsonb('metadata').$type<MediaMetadata>().default({}),
  ...timestamps(),
  ...softDelete(),
}, (table) => [
  index('idx_media_tenant').on(table.tenantId),
  index('idx_media_module').on(table.moduleId),
  index('idx_media_type').on(table.type),
  index('idx_media_transcription_status').on(table.transcriptionStatus),
  pgPolicy('media_tenant_isolation', {

```

```

for: 'all', to: appRole,
using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
]);

export interface MediaMetadata {
width?: number;
height?: number;
codec?: string;
bitrate?: number;
language?: string;
processingJobId?: string;
}

// -----
// transcripts — Full transcription per media asset
// -----

export const transcripts = pgTable('transcripts', {
id: pk(),
tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
assetId: pk().references(() => mediaAssets.id, { onDelete: 'cascade' }).notNull(),
fullText: text('full_text').notNull(),
language: varchar('language', { length: 10 }).notNull().default('he'),
modelUsed: varchar('model_used', { length: 100 }), // e.g., 'whisper-large-v3'
confidence: real('confidence'), // average confidence 0-1
wordCount: integer('word_count'),
...timestamps(),
}, (table) => [
index('idx_transcripts_tenant').on(table.tenantId),
index('idx_transcripts_asset').on(table.assetId),

pgPolicy('transcripts_tenant_isolation', {
for: 'all', to: appRole,
using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
]);

// -----
// transcript_segments — Time-aligned segments
// -----

export const transcriptSegments = pgTable('transcript_segments', {
id: pk(),
tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
transcriptId: pk().references(() => transcripts.id, { onDelete: 'cascade' }).notNull(),
assetId: pk().references(() => mediaAssets.id, { onDelete: 'cascade' }).notNull(),
startTime: real('start_time').notNull(), // seconds from beginning
endTime: real('end_time').notNull(), // seconds from beginning
});

```

```
text: text('text').notNull(),
speaker: varchar('speaker', { length: 255 }),      // speaker diarization
confidence: real('confidence'),
...timestamps(),
}, (table) => [
  index('idx_segments_tenant').on(table.tenantId),
  index('idx_segments_transcript').on(table.transcriptId),
  index('idx_segments_asset').on(table.assetId),
  index('idx_segments_time_range').on(table.assetId, table.startTime, table.endTime),

  pgPolicy('segments_tenant_isolation', {
    for: 'all', to: appRole,
    using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
  }),
];

```

6. Annotation Schema

typescript

```

// File: packages/db/src/schema/annotations.ts

import {
  pgTable, pgPolicy, varchar, text, jsonb,
  real, integer, index, uuid as uuidCol,
} from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {

  pk, tenantId, timestamps, softDelete,
  annotationTypeEnum, annotationLayerEnum,
  appRole, rlsTenantCheck, rlsTenantInsertCheck,
} from './_shared.ts';
import { tenants } from './tenants.ts';
import { users } from './users.ts';
import { mediaAssets } from './media-assets.ts';

// -----
// annotations — All user/AI markings on content
// -----

export const annotations = pgTable('annotations', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  userId: uuidCol('user_id').notNull().references(() => users.id, { onDelete: 'cascade' }),
  assetId: uuidCol('asset_id').notNull().references(() => mediaAssets.id, { onDelete: 'cascade' }),
  type: annotationTypeEnum('type').notNull(),
  layer: annotationLayerEnum('layer').notNull().default('personal'),

  // Content — depends on type
  content: text('content'),           // text content or markdown
  sketchData: jsonb('sketch_data').$type<SketchData>(), // Konva.js JSON for sketches

  // Temporal anchoring
  timestampStart: real('timestamp_start'),    // seconds (null for non-temporal)
  timestampEnd: real('timestamp_end'),        // seconds (null for point annotations)

  // Spatial anchoring (for spatial_comment on video frame)
  spatialX: real('spatial_x'),           // 0-1 normalized X position
  spatialY: real('spatial_y'),           // 0-1 normalized Y position

  // PDF-specific anchoring
  pageNumber: integer('page_number'),
  textRangeStart: integer('text_range_start'), // character offset
  textRangeEnd: integer('text_range_end'),   // character offset

  // Reference to other annotation (reply chain)
  parentId: uuidCol('parent_id').references((): any => annotations.id, { onDelete: 'cascade' }),
}

```

```

// AI-generated metadata
agentId: uuidCol('agent_id'),           // if created by an agent
metadata: jsonb('metadata').$type<AnnotationMetadata>().default({}),
          ...timestamps(),
          ...softDelete(),
}, (table) => [
  index('idx_annotations_tenant').on(table.tenantId),
  index('idx_annotations_user').on(table.userId),
  index('idx_annotations_asset').on(table.assetId),
  index('idx_annotations_asset_time').on(table.assetId, table.timestampStart),
  index('idx_annotations_layer').on(table.layer),
  index('idx_annotations_parent').on(table.parentId),
}

pgPolicy('annotations_select_tenant', {
  for: 'select', to: appRole, using: rlsTenantCheck,
}),
pgPolicy('annotations_insert_tenant', {
  for: 'insert', to: appRole, withCheck: rlsTenantInsertCheck,
}),
pgPolicy('annotations_update_own', {
  for: 'update', to: appRole,
  // Users can only update their own annotations within same tenant
  using: sql`tenant_id = current_setting('app.current_tenant', true)::uuid
    AND user_id = current_setting('app.current_user', true)::uuid`,
  withCheck: rlsTenantInsertCheck,
}),
pgPolicy('annotations_delete_own', {
  for: 'delete', to: appRole,
  using: sql`tenant_id = current_setting('app.current_tenant', true)::uuid
    AND user_id = current_setting('app.current_user', true)::uuid`,
}),
);

```

```

export interface SketchData {
  /** Konva.js stage JSON export */
  stage: Record<string, unknown>;
  /** Dimensions the sketch was created at */
  canvasWidth: number;
  canvasHeight: number;
  /** Video frame timestamp this sketch is anchored to */
  frameTimestamp?: number;
}

```

```

export interface AnnotationMetadata {
  color?: string;
}

```

```
pinned?: boolean;  
resolvedAt?: string;  
linkedConceptIds?: string[];  
aiConfidence?: number;  
}
```

7. Collaboration Schema — CRDT Persistence

typescript

```
// File: packages/db/src/schema/collaboration.ts

import {
  pgTable, pgPolicy, varchar, text, jsonb,
  integer, index, uuid as uuidCol, customType,
} from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {

  pk, tenantId, timestamps,
  appRole, rlsTenantCheck, rlsTenantInsertCheck,
} from './_shared.ts';
import { tenants } from './tenants.ts';
import { users } from './users.ts';

// -----
// Custom type: bytea for CRDT binary data
// -----

const bytea = customType<{ data: Buffer; driverData: Buffer }>({
  dataType() {
    return 'bytea';
  },
});

// -----
// collab_documents — Root document for Yjs collaboration
// -----

export const collabDocuments = pgTable('collab_documents', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  /** Hocuspocus document name — unique per tenant */
  documentName: varchar('document_name', { length: 500 }).notNull(),
  /** The compacted Yjs state vector (full snapshot) */
  stateVector: bytea('state_vector'),
  /** Number of incremental updates since last compaction */
  pendingUpdates: integer('pending_updates').notNull().default(0),
  /** Human-readable title */
  title: varchar('title', { length: 500 }),
  /** Associated entity (annotation layer, course notes, etc.) */
  entityType: varchar('entity_type', { length: 100 }),
  entityId: uuidCol('entity_id'),
  ...timestamps(),
}, (table) => [
  index('idx_collab_docs_tenant').on(table.tenantId),
  index('idx_collab_docs_name').on(table.tenantId, table.documentName),
  index('idx_collab_docs_entity').on(table.entityType, table.entityId),
]
```

```

pgPolicy('collab_docs_tenant_isolation', {
  for: 'all', to: appRole,
  using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
);

// -----
// crdt_updates — Incremental Yjs binary updates
// -----

// Append-only. Periodically compacted into collab_documents.state_vector
export const crdtUpdates = pgTable('crdt_updates', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  documentId: uuidCol('document_id').notNull()
    .references(() => collabDocuments.id, { onDelete: 'cascade' }),
  /** The binary Yjs update */
  update: bytea('update').notNull(),
  /** Sequence number for ordering */
  seq: integer('seq').notNull(),
  /** Which user caused this update */
  userId: uuidCol('user_id').references(() => users.id, { onDelete: 'set null' }),
  ...timestamps(),
}, (table) => [
  index('idx_crdt_updates_doc').on(table.documentId),
  index('idx_crdt_updates_doc_seq').on(table.documentId, table.seq),
]

pgPolicy('crdt_updates_tenant_isolation', {
  for: 'all', to: appRole,
  using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
);

// -----
// collab_sessions — Active collaboration sessions (ephemeral)
// -----

export const collabSessions = pgTable('collab_sessions', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  documentId: uuidCol('document_id').notNull()
    .references(() => collabDocuments.id, { onDelete: 'cascade' }),
  userId: uuidCol('user_id').notNull().references(() => users.id, { onDelete: 'cascade' }),
  cursorPosition: jsonb('cursor_position').$type<CursorPosition>(),
  connectedAt: timestamps().createdAt,
  lastSeenAt: timestamps().updatedAt,
}, (table) => [
  index('idx_collab_sessions_doc').on(table.documentId),
]

```

```
pgPolicy('collab_sessions_tenant_isolation', {
  for: 'all', to: appRole,
  using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
]);



```

```
export interface CursorPosition {
  x?: number;
  y?: number;
  timestamp?: number; // video timestamp if applicable
  pageNumber?: number; // PDF page if applicable
  userName?: string;
  color?: string;
}
```

8. Agent Schema

typescript

```
// File: packages/db/src/schema/agents.ts

import {
  pgTable, pgPolicy, varchar, text, jsonb,
  integer, real, index, uuid as uuidCol,
  boolean,
} from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {

  pk, tenantId, timestamps, softDelete,
  agentTemplateEnum, agentTriggerEnum, agentOutputFormatEnum,
  executionStatusEnum,
  appRole, rlsTenantCheck, rlsTenantInsertCheck,
} from './_shared.ts';
import { tenants } from './tenants.ts';
import { users } from './users.ts';

//
```

```
// _____
// agent_definitions — User-created or built-in agents
// _____
export const agentDefinitions = pgTable('agent_definitions', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  creatorId: uuidCol('creator_id').notNull().references(() => users.id, { onDelete: 'cascade' }),
  name: varchar('name', { length: 255 }).notNull(),
  description: text('description'),
  templateType: agentTemplateEnum('template_type').notNull(),
  isBuiltIn: boolean('is_built_in').notNull().default(false),
  isPublic: boolean('is_public').notNull().default(false),

  // Agent configuration
  config: jsonb('config').$type<AgentConfig>().notNull(),

  // Usage stats
  executionCount: integer('execution_count').notNull().default(0),
  avgExecutionMs: real('avg_execution_ms'),
  lastExecutedAt: timestamps().updatedAt,
  ...timestamps(),
  ...softDelete(),
}, (table) => [
  index('idx_agent_defs_tenant').on(table.tenantId),
  index('idx_agent_defs_creator').on(table.creatorId),
  index('idx_agent_defs_template').on(table.templateType),
  pgPolicy('agent_defs_select_tenant', {
```

```

    for: 'select', to: appRole, using: rlsTenantCheck,
}),
pgPolicy('agent_defs_insert_tenant', {
    for: 'insert', to: appRole, withCheck: rlsTenantInsertCheck,
}),
pgPolicy('agent_defs_update_own', {
    for: 'update', to: appRole,
    using: sql`tenant_id = current_setting('app.current_tenant', true)::uuid
        AND creator_id = current_setting('app.current_user', true)::uuid`,
    withCheck: rlsTenantInsertCheck,
}),
pgPolicy('agent_defs_delete_own', {
    for: 'delete', to: appRole,
    using: sql`tenant_id = current_setting('app.current_tenant', true)::uuid
        AND creator_id = current_setting('app.current_user', true)::uuid`,
}),
]);

```

```

export interface AgentConfig {
    dataScope: 'MY_NOTES' | 'COURSE' | 'ORGANIZATION' | 'CUSTOM_QUERY';
    customQuery?: string;           // Cypher query for CUSTOM_QUERY scope
    triggerType: AgentTrigger;
    outputFormat: AgentOutputFormat;
    personality?: string;          // e.g., "Strict Talmudic scholar"
    difficulty?: 'BEGINNER' | 'INTERMEDIATE' | 'ADVANCED';
    modelOverride?: string;         // e.g., 'phi4:14b'
    maxTokens?: number;
    temperature?: number;
    systemPrompt?: string;
    toolsEnabled?: string[];        // which MCP tools the agent can use
}

```

```

// -----
// agent_executions — Log of every agent run
// -----
export const agentExecutions = pgTable('agent_executions', {
    id: pk(),
    tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
    agentId: uuidCol('agent_id').notNull()
        .references(() => agentDefinitions.id, { onDelete: 'cascade' }),
    userId: uuidCol('user_id').notNull().references(() => users.id, { onDelete: 'cascade' }),
    status: executionStatusEnum('status').notNull().default('queued'),
    // Input/Output
    input: jsonb('input').$type<Record<string, unknown>>().notNull(),
    output: jsonb('output').$type<Record<string, unknown>>(),
    reasoning: text('reasoning'),      // chain-of-thought (debugging)
})
```

```
errorMessage: text('error_message'),  
  
// Performance tracking  
tokensUsed: integer('tokens_used'),  
executionMs: integer('execution_ms'),  
modelUsed: varchar('model_used', { length: 100 }),  
  
...timestamps(),  
, (table) => [  
    index('idx_agent_execs_tenant').on(table.tenantId),  
    index('idx_agent_execs_agent').on(table.agentId),  
    index('idx_agent_execs_user').on(table.userId),  
    index('idx_agent_execs_status').on(table.status),  
  
    pgPolicy('agent_execs_tenant_isolation', {  
        for: 'all', to: appRole,  
        using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,  
    }),  
];
```

9. Row-Level Security Policies

typescript

```

// File: packages/db/src/schema/_rls-policies.ts
//
// This file documents the RLS strategy. Actual policies are defined
// inline on each table (see above). This file provides the middleware
// that sets the tenant context per request.

/***
 * CRITICAL: Every database transaction MUST set the tenant context.
 * This is enforced by the `withTenantContext` wrapper.
 *
 * Pattern:
 *  SET LOCAL app.current_tenant = '<tenant-uuid>';
 *  SET LOCAL app.current_user = '<user-uuid>';
 *  SET LOCAL ROLE edusphere_app;
 *
 * SET LOCAL is transaction-scoped — auto-resets on commit/rollback.
 * This is safe with PgBouncer in transaction mode.
 */

```

```

import { sql } from 'drizzle-orm';
import { type PostgresJsDatabase } from 'drizzle-orm/postgres-js';

```

```

export interface TenantContext {
  tenantId: string;
  userId: string;
}

```

```

/***
 * Execute a callback within a tenant-isolated transaction.
 * ALL user-facing database operations MUST use this wrapper.
 *
 * @example
 * ```typescript
 * const result = await withTenantContext(db, ctx, async (tx) => {
 *   return tx.select().from(courses);
 *   // RLS automatically filters by tenant_id
 * });
 * ```
 */

```

```

export async function withTenantContext<T>(
  db: PostgresJsDatabase,
  ctx: TenantContext,
  callback: (tx: PostgresJsDatabase) => Promise<T>,
): Promise<T> {
  return db.transaction(async (tx) => {
    // Set tenant context for RLS
  })
}

```

```

await tx.execute(
  sql`SET LOCAL app.current_tenant = ${ctx.tenantId}`
);
await tx.execute(
  sql`SET LOCAL app.current_user = ${ctx.userId}`
);
// Switch to application role (RLS policies apply)
await tx.execute(sql`SET LOCAL ROLE edusphere_app`);

return callback(tx);
});

}

/***
 * For service-to-service calls that need to bypass RLS
 * (e.g., NATS event handlers processing cross-tenant data).
 * Use SPARINGLY and document why.
*/
export async function withSuperuserContext<T>(
  db: PostgresJsDatabase,
  callback: (tx: PostgresJsDatabase) => Promise<T>,
): Promise<T> {
  // No SET LOCAL ROLE — runs as connection owner (superuser)
  return db.transaction(async (tx) => callback(tx));
}

```

RLS Policy Summary Table:

Table	SELECT	INSERT	UPDATE	DELETE
tenants	Own tenant only	Super admin	Own tenant	Super admin
users	Same tenant	Same tenant	Same tenant	—
courses	Same tenant	Same tenant	Same tenant	Same tenant
modules	Same tenant	Same tenant	Same tenant	Same tenant
media_assets	Same tenant	Same tenant	Same tenant	Same tenant
transcripts	Same tenant	Same tenant	Same tenant	Same tenant
transcript_segments	Same tenant	Same tenant	Same tenant	Same tenant
annotations	Same tenant	Same tenant	Own only + tenant	Own only + tenant
collab_documents	Same tenant	Same tenant	Same tenant	Same tenant
crdt_updates	Same tenant	Same tenant	Same tenant	Same tenant
agent_definitions	Same tenant	Same tenant	Own only + tenant	Own only + tenant
agent_executions	Same tenant	Same tenant	Same tenant	Same tenant

10. Apache AGE — Knowledge Graph Ontology

10.1 Graph Initialization

sql

```

-- File: packages/db/src/graph/init-graph.sql
-- Run after extensions are created

-- Create vertex labels (node types)
SELECT create_vlabel('eduspHERE_graph', 'Concept');
SELECT create_vlabel('eduspHERE_graph', 'Person');
SELECT create_vlabel('eduspHERE_graph', 'Term');
SELECT create_vlabel('eduspHERE_graph', 'Source');
SELECT create_vlabel('eduspHERE_graph', 'TopicCluster');

-- Create edge labels (relationship types)
SELECT create_elabel('eduspHERE_graph', 'MENTIONS');
SELECT create_elabel('eduspHERE_graph', 'REFERS_TO');
SELECT create_elabel('eduspHERE_graph', 'RELATED_TO');
SELECT create_elabel('eduspHERE_graph', 'CONTRADICTS');
SELECT create_elabel('eduspHERE_graph', 'PREREQUISITE_OF');
SELECT create_elabel('eduspHERE_graph', 'DERIVED_FROM');
SELECT create_elabel('eduspHERE_graph', 'AUTHORED_BY');
SELECT create_elabel('eduspHERE_graph', 'CITES');
SELECT create_elabel('eduspHERE_graph', 'INFERRRED_RELATED');
SELECT create_elabel('eduspHERE_graph', 'BELONGS_TO');

-- Create indexes on vertex properties for fast lookups
CREATE INDEX idx_concept_tenant ON eduspHERE_graph."Concept" USING btree ((properties->>'tenant_id'));
CREATE INDEX idx_concept_label ON eduspHERE_graph."Concept" USING gin ((properties->'label'));
CREATE INDEX idx_person_name ON eduspHERE_graph."Person" USING gin ((properties->'name'));
CREATE INDEX idx_term_label ON eduspHERE_graph."Term" USING gin ((properties->'label'));

```

10.2 Vertex Property Schemas

typescript

```

// File: packages/db/src/graph/ontology.ts

/**
 * Vertex (Node) property schemas for the knowledge graph.
 * These are stored as agtype JSON in Apache AGE.
 * Validated with Zod before insertion.
 */

import { z } from 'zod';

// — Concept ——————
// A semantic concept (e.g., "Thermodynamics", "Rambam", "Polymers")
export const ConceptVertex = z.object({
  id: z.string().uuid(), // Same UUID as concepts SQL table
  tenant_id: z.string().uuid(),
  label: z.string().min(1), // Primary label (Hebrew/English)
  description: z.string().optional(),
  aliases: z.array(z.string()).default([]), // Synonyms, abbreviations
  domain: z.string().optional(), // e.g., "Physics", "Talmud", "Engineering"
  confidence: z.number().min(0).max(1).default(1), // 1 = manual, <1 = AI-extracted
  source_type: z.enum(['manual', 'ai_extracted', 'imported']).default('manual'),
  created_by: z.string().uuid().optional(),
});

export type ConceptVertex = z.infer<typeof ConceptVertex>;

// — Person ——————
// A person referenced in content (author, speaker, historical figure)
export const PersonVertex = z.object({
  id: z.string().uuid(),
  tenant_id: z.string().uuid(),
  name: z.string(),
  role: z.string().optional(), // e.g., "Speaker", "Author", "Historical"
  external_url: z.string().url().optional(),
});

export type PersonVertex = z.infer<typeof PersonVertex>;

// — Term ——————
// A domain-specific term (narrower than Concept)
export const TermVertex = z.object({
  id: z.string().uuid(),
  tenant_id: z.string().uuid(),
  label: z.string(),
  domain: z.string(),
  definition: z.string().optional(),
});

export type TermVertex = z.infer<typeof TermVertex>;

```

```

// — Source ——————
// An external reference (book, paper, URL)
export const SourceVertex = z.object({
  id: z.string().uuid(),
  tenant_id: z.string().uuid(),
  title: z.string(),
  type: z.enum(['book', 'paper', 'url', 'lecture', 'internal_asset']),
  external_url: z.string().optional(),
  internal_asset_id: z.string().uuid().optional(), // FK to media_assets
  isbn: z.string().optional(),
  doi: z.string().optional(),
});
export type SourceVertex = z.infer<typeof SourceVertex>

// — TopicCluster ——————
// Auto-generated grouping of related concepts
export const TopicClusterVertex = z.object({
  id: z.string().uuid(),
  tenant_id: z.string().uuid(),
  label: z.string(),
  concept_count: z.number().int(),
  coherence_score: z.number().min(0).max(1),
});
export type TopicClusterVertex = z.infer<typeof TopicClusterVertex>;

```

10.3 Edge Property Schemas

typescript

```
// File: packages/db/src/graph/ontology.ts (continued)
```

```
// — Edge schemas ——————
```

```
export const MentionsEdge = z.object({
  segment_id: z.string().uuid(), // FK to transcript_segments
  asset_id: z.string().uuid(), // FK to media_assets
  start_time: z.number(),
  end_time: z.number(),
  confidence: z.number().min(0).max(1).default(1),
});
```

```
export const RelatedToEdge = z.object({
  strength: z.number().min(0).max(1), // Semantic similarity score
  method: z.enum(['manual', 'embedding_similarity', 'co_occurrence', 'citation']),
  evidence: z.string().optional(), // Human-readable explanation
});
```

```
export const ContradictsEdge = z.object({
  evidence: z.string(), // Why these concepts contradict
  source_a: z.string().uuid(), // Which media asset/annotation
  source_b: z.string().uuid(),
  severity: z.enum(['minor', 'moderate', 'major']).default('moderate'),
  detected_by: z.enum(['manual', 'chavruta_agent', 'ai_inference']),
});
```

```
export const PrerequisiteOfEdge = z.object({
  strength: z.enum(['required', 'recommended', 'optional']).default('recommended'),
});
```

```
export const InferredRelatedEdge = z.object({
  confidence: z.number().min(0).max(1),
  method: z.string(), // Algorithm used
  inferred_at: z.string().datetime(),
  reviewed: z.boolean().default(false),
});
```

```
export const CitesEdge = z.object({
  page: z.number().optional(),
  timestamp: z.number().optional(), // If citing a video moment
  context: z.string().optional(), // Quote or surrounding text
});
```

10.4 Example Cypher Queries

```
// File: packages/db/src/graph/queries.ts
```

```
import { sql } from 'drizzle-orm';
import { GRAPH_NAME } from './schema/_shared.ts';

/***
 * All Cypher queries are executed via SQL using Apache AGE's cypher() function.
 * Parameters are passed as the third argument (agtype) for safety.
 */

```

```
/** Create a new Concept vertex */
export const createConceptQuery = (concept: ConceptVertex) => sql`  
SELECT * FROM cypher(${GRAPH_NAME}, $$  
CREATE (c:Concept {  
    id: '${sql.raw(concept.id)}',  
    tenant_id: '${sql.raw(concept.tenant_id)}',  
    label: '${sql.raw(concept.label)}',  
    aliases: ${sql.raw(JSON.stringify(concept.aliases))},  
    domain: '${sql.raw(concept.domain ?? "")}',  
    confidence: ${sql.raw(String(concept.confidence))},  
    source_type: '${sql.raw(concept.source_type)}'  
})  
RETURN c  
$$) AS (v agtype)  
`;
```

```
/** Find all concepts related to a specific concept */
export const findRelatedConceptsQuery = (conceptId: string, tenantId: string) => sql`  
SELECT * FROM cypher(${GRAPH_NAME}, $$  
MATCH (c:Concept {id: '${sql.raw(conceptId)}', tenant_id: '${sql.raw(tenantId)}'})  
    -[r:RELATED_TO|CONTRADICTS|PREREQUISITE_OF*1..3]-  
    (related:Concept)  
WHERE related.tenant_id = '${sql.raw(tenantId)}'  
RETURN related.id, related.label, type(r) AS relation_type, r.strength  
ORDER BY r.strength DESC  
LIMIT 50  
$$) AS (id agtype, label agtype, relation_type agtype, strength agtype)  
`;
```

```
/** Get semantic context for a video timestamp (for Contextual Copilot) */
export const getSemanticContextQuery = (assetId: string, timestamp: number, tenantId: string) => sql`  
SELECT * FROM cypher(${GRAPH_NAME}, $$  
MATCH (seg)-[:MENTIONS]->(concept:Concept)  
WHERE seg.asset_id = '${sql.raw(assetId)}'  
    AND concept.tenant_id = '${sql.raw(tenantId)}'  
    AND seg.start_time <= ${sql.raw(String(timestamp))}
```

```

        AND seg.end_time >= ${sql.raw(String(timestamp))}

WITH concept
OPTIONAL MATCH (concept)-[r:RELATED_TO|CONTRADICTS]->(related:Concept)
WHERE related.tenant_id = '${sql.raw(tenantId)}'
RETURN concept.id, concept.label, concept.domain,
       collect({
         related_id: related.id,
         related_label: related.label,
         relation: type(r),
         strength: r.strength
       }) AS connections
$$) AS (id agtype, label agtype, domain agtype, connections agtype)
';

/** Find contradictions for Chavruta agent */
export const findContradictionsQuery = (conceptId: string, tenantId: string) => sql` 
SELECT * FROM cypher(${GRAPH_NAME}, $$ 
  MATCH (a:Concept {id: '${sql.raw(conceptId)}', tenant_id: '${sql.raw(tenantId)}'}) 
    -[r:CONTRADICTS]- 
    (b:Concept {tenant_id: '${sql.raw(tenantId)}'}) 
  RETURN b.id, b.label, r.evidence, r.severity, r.detected_by
$$) AS (id agtype, label agtype, evidence agtype, severity agtype, detected_by agtype)
`;

/** Build concept prerequisite chain (learning path) */
export const getPrerequisiteChainQuery = (conceptId: string, tenantId: string) => sql` 
SELECT * FROM cypher(${GRAPH_NAME}, $$ 
  MATCH path = (target:Concept {id: '${sql.raw(conceptId)}', tenant_id: '${sql.raw(tenantId)}'}) 
    -[:PREREQUISITE_OF*1..5]-
    (prereq:Concept {tenant_id: '${sql.raw(tenantId)}'}) 
  RETURN nodes(path) AS chain, length(path) AS depth
  ORDER BY depth ASC
$$) AS (chain agtype, depth agtype)
`;

```

11. pgvector — Embedding Tables & Indexes

typescript

```

// File: packages/db/src/schema/embeddings.ts

import {
  pgTable, pgPolicy, varchar, text, jsonb,
  integer, index, uuid as uuidCol, real,
} from 'drizzle-orm/pg-core';
import { sql } from 'drizzle-orm';
import {

  pk, tenantId, timestamps,
  appRole, rlsTenantCheck, rlsTenantInsertCheck,
  EMBEDDING_DIMENSIONS,
} from './_shared.ts';
import { tenants } from './tenants.ts';
import { mediaAssets } from './media-assets.ts';
import { transcriptSegments } from './media-assets.ts';
import { annotations } from './annotations.ts';


$$/*
 * Custom vector column type for pgvector.
 * Drizzle does not have built-in vector support,
 * so we define a custom type.
 */

import { customType } from 'drizzle-orm/pg-core';

const vector = (name: string, dimensions: number) =>
  customType<{ data: number[]; driverData: string }>({
    dataType() {
      return `vector(${dimensions})`;
    },
    toDriver(value: number[]) {
      return `[${value.join(',')}]`;
    },
    fromDriver(value: string) {
      return value
        .replace('T', " ")
        .replace('J', " ")
        .split(',')
        .map(Number);
    },
  })(name);

// -----
// content_embeddings — Vectors for transcript segments
// -----$$


export const contentEmbeddings = pgTable('content_embeddings', {
  id: pk(),

```

```

tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
segmentId: uuidCol('segment_id')

.references(() => transcriptSegments.id, { onDelete: 'cascade' }),
assetId: uuidCol('asset_id').notNull()

.references(() => mediaAssets.id, { onDelete: 'cascade' }),

/** The text that was embedded */

sourceText: text('source_text').notNull(),

/** The embedding vector */

embedding: vector('embedding', EMBEDDING_DIMENSIONS).notNull(),

/** Which model generated this embedding */

modelName: varchar('model_name', { length: 100 }).notNull(),

/** Chunk index if text was split */

chunkIndex: integer('chunk_index').notNull().default(0),

...timestamps(),

}, (table) => [
  index('idx_content_emb_tenant').on(table.tenantId),
  index('idx_content_emb_asset').on(table.assetId),
]

pgPolicy('content_embeddings_tenant_isolation', {
  for: 'all', to: appRole,
  using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
}),
);

//
```

```

// annotation_embeddings — Vectors for user annotations
//
```

```

export const annotationEmbeddings = pgTable('annotation_embeddings', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  annotationId: uuidCol('annotation_id').notNull()

  .references(() => annotations.id, { onDelete: 'cascade' }),

  sourceText: text('source_text').notNull(),
  embedding: vector('embedding', EMBEDDING_DIMENSIONS).notNull(),
  modelName: varchar('model_name', { length: 100 }).notNull(),
  ...timestamps(),

}, (table) => [
  index('idx_annotation_emb_tenant').on(table.tenantId),

  pgPolicy('annotation_embeddings_tenant_isolation', {
    for: 'all', to: appRole,
    using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
  }),
);

//
```

```

// concept_embeddings — Vectors for knowledge graph concepts

```

```
// -----
export const conceptEmbeddings = pgTable('concept_embeddings', {
  id: pk(),
  tenantId: tenantId().references(() => tenants.id, { onDelete: 'restrict' }),
  /** UUID matching the Concept vertex in AGE */
  conceptId: uuidCol('concept_id').notNull(),
  /** label + description + aliases concatenated */
  sourceText: text('source_text').notNull(),
  embedding: vector('embedding', EMBEDDING_DIMENSIONS).notNull(),
  modelName: varchar('model_name', { length: 100 }).notNull(),
  ...timestamps(),
}, (table) => [
  index('idx_concept_emb_tenant').on(table.tenantId),
  index('idx_concept_emb_concept').on(table.conceptId),

  pgPolicy('concept_embeddings_tenant_isolation', {
    for: 'all', to: appRole,
    using: rlsTenantCheck, withCheck: rlsTenantInsertCheck,
  }),
]);

```

11.1 HNSW Indexes (Created via Migration)

```
sql
-- File: packages/db/src/migrations/0002_vector_indexes.sql
-- HNSW indexes for approximate nearest neighbor search
-- cosine distance is best for normalized embeddings (sentence-transformers, nomic-embed-text)

-- Content embeddings index
CREATE INDEX CONCURRENTLY idx_content_emb_hnsw
ON content_embeddings
USING hnsw (embedding vector_cosine_ops)
WITH (m = 16, ef_construction = 200);

-- Annotation embeddings index
CREATE INDEX CONCURRENTLY idx_annotation_emb_hnsw
ON annotation_embeddings
USING hnsw (embedding vector_cosine_ops)
WITH (m = 16, ef_construction = 200);

-- Concept embeddings index
CREATE INDEX CONCURRENTLY idx_concept_emb_hnsw
ON concept_embeddings
USING hnsw (embedding vector_cosine_ops)
WITH (m = 16, ef_construction = 200);
```

11.2 Semantic Search Functions

sql

```

-- File: packages/db/src/migrations/0003_search_functions.sql

-- Semantic search across all content within a tenant
CREATE OR REPLACE FUNCTION search_content_semantic(
    query_embedding vector(768),
    p_tenant_id uuid,
    p_limit integer DEFAULT 10,
    p_min_similarity float DEFAULT 0.7
)
RETURNS TABLE (
    segment_id uuid,
    asset_id uuid,
    source_text text,
    similarity float
)
LANGUAGE sql STABLE
AS $$

SELECT
    ce.segment_id,
    ce.asset_id,
    ce.source_text,
    1 - (ce.embedding <=> query_embedding) AS similarity
FROM content_embeddings ce
WHERE ce.tenant_id = p_tenant_id
    AND 1 - (ce.embedding <=> query_embedding) >= p_min_similarity
ORDER BY ce.embedding <=> query_embedding
LIMIT p_limit;

$$;

-- HybridRAG: combines vector search + graph traversal
-- Called from the semantic service, results merged in TypeScript
CREATE OR REPLACE FUNCTION search_concepts_semantic(
    query_embedding vector(768),
    p_tenant_id uuid,
    p_limit integer DEFAULT 10
)
RETURNS TABLE (
    concept_id uuid,
    source_text text,
    similarity float
)
LANGUAGE sql STABLE
AS $$

SELECT
    ce.concept_id,
    ce.source_text,

```

```
1 - (ce.embedding <=> query_embedding) AS similarity
FROM concept_embeddings ce
WHERE ce.tenant_id = p_tenant_id
ORDER BY ce.embedding <=> query_embedding
LIMIT p_limit;
$$;
```

12. Drizzle Relations (RQBv2)

typescript

```
// File: packages/db/src/relations.ts

import { defineRelations } from 'drizzle-orm';
import * as schema from './schema/index.ts';

export const relations = defineRelations(schema, (r) => ({
  // — Tenants ——————
  tenants: {
    users: r.many.users({ from: r.tenants.id, to: r.users.tenantId }),
    courses: r.many.courses({ from: r.tenants.id, to: r.courses.tenantId }),
  },
  // — Users ——————
  users: {
    tenant: r.one.tenants({ from: r.users.tenantId, to: r.tenants.id }),
    annotations: r.many.annotations({ from: r.users.id, to: r.annotations.userId }),
    agentDefinitions: r.many.agentDefinitions({ from: r.users.id, to: r.agentDefinitions.creatorId }),
    agentExecutions: r.many.agentExecutions({ from: r.users.id, to: r.agentExecutions.userId }),
  },
  // — Courses ——————
  courses: {
    tenant: r.one.tenants({ from: r.courses.tenantId, to: r.tenants.id }),
    modules: r.many.modules({ from: r.courses.id, to: r.modules.courseId }),
    forkedFrom: r.one.courses({ from: r.courses.forkedFromId, to: r.courses.id }),
  },
  // — Modules ——————
  modules: {
    course: r.one.courses({ from: r.modules.courseId, to: r.courses.id }),
    mediaAssets: r.many.mediaAssets({ from: r.modules.id, to: r.mediaAssets.moduleId }),
  },
  // — Media Assets ——————
  mediaAssets: {
    module: r.one.modules({ from: r.mediaAssets.moduleId, to: r.modules.id }),
    transcripts: r.many.transcripts({ from: r.mediaAssets.id, to: r.transcripts.assetId }),
    annotations: r.many.annotations({ from: r.mediaAssets.id, to: r.annotations.assetId }),
    contentEmbeddings: r.many.contentEmbeddings({ from: r.mediaAssets.id, to: r.contentEmbeddings.assetId }),
  },
  // — Transcripts ——————
  transcripts: {
    asset: r.one.mediaAssets({ from: r.transcripts.assetId, to: r.mediaAssets.id }),
    segments: r.many.transcriptSegments({ from: r.transcripts.id, to: r.transcriptSegments.transcriptId }),
  },
})
```

```

// — Transcript Segments ——————
transcriptSegments: {
  transcript: r.one.transcripts({ from: r.transcriptSegments.transcriptId, to: r.transcripts.id }),
  asset: r.one.mediaAssets({ from: r.transcriptSegments.assetId, to: r.mediaAssets.id }),
},
}

// — Annotations ——————
annotations: {
  user: r.one.users({ from: r.annotations.userId, to: r.users.id }),
  asset: r.one.mediaAssets({ from: r.annotations.assetId, to: r.mediaAssets.id }),
  parent: r.one.annotations({ from: r.annotations.parentId, to: r.annotations.id }),
  replies: r.many.annotations({ from: r.annotations.id, to: r.annotations.parentId }),
},
}

// — Agent Definitions ——————
agentDefinitions: {
  creator: r.one.users({ from: r.agentDefinitions.creatorId, to: r.users.id }),
  executions: r.many.agentExecutions({ from: r.agentDefinitions.id, to: r.agentExecutions.agentId }),
},
}

// — Agent Executions ——————
agentExecutions: {
  agent: r.one.agentDefinitions({ from: r.agentExecutions.agentId, to: r.agentDefinitions.id }),
  user: r.one.users({ from: r.agentExecutions.userId, to: r.users.id }),
},
}

// — Collab Documents ——————
collabDocuments: {
  updates: r.many.crdtUpdates({ from: r.collabDocuments.id, to: r.crdtUpdates.documentId }),
  sessions: r.many.collabSessions({ from: r.collabDocuments.id, to: r.collabSessions.documentId }),
},
});

}

```

13. Migration Strategy

13.1 Drizzle Kit Configuration

typescript

```
// File: packages/db/drizzle.config.ts

import { defineConfig } from 'drizzle-kit';

export default defineConfig({
  schema: './src/schema/index.ts',
  out: './src/migrations',
  dialect: 'postgresql',
  dbCredentials: {
    url: process.env.DATABASE_URL!,
  },
  verbose: true,
  strict: true,
  // Manage RLS policies and roles
  entities: {
    roles: {
      provider: "", // Empty = drizzle manages roles
      exclude: ['postgres', 'pg_*'],
    },
  },
});
});
```

13.2 Migration Order

1. 0000_extensions.sql — PG extensions (manual, run as superuser)
2. 0001_initial_schema.sql — Drizzle-generated tables + RLS policies
3. 0002_vector_indexes.sql — HNSW indexes (manual, CONCURRENTLY)
4. 0003_search_functions.sql — Semantic search PG functions (manual)
5. 0004_graph_init.sql — Apache AGE graph + labels + indexes (manual)

13.3 Commands

```
bash

# Generate migration from schema changes
pnpm --filter @edusphere/db exec drizzle-kit generate

# Apply migrations
pnpm --filter @edusphere/db exec drizzle-kit migrate

# Push schema directly (dev only)
pnpm --filter @edusphere/db exec drizzle-kit push

# Open Drizzle Studio (visual DB browser)
pnpm --filter @edusphere/db exec drizzle-kit studio
```

14. Seed Data

```
typescript
```

```
// File: packages/db/src/seed/index.ts

import { db } from './client.ts';
import { tenants, users, courses, modules, mediaAssets } from './schema/index.ts';

const DEFAULT_TENANT_ID = '00000000-0000-0000-0000-000000000001';
const ADMIN_USER_ID = '00000000-0000-0000-0000-000000000010';
const STUDENT_USER_ID = '00000000-0000-0000-0000-000000000020';

export async function seed() {
  console.log('🌱 Seeding database...');

  // 1. Create default tenant
  await db.insert(tenants).values({
    id: DEFAULT_TENANT_ID,
    name: 'EduSphere Dev',
    slug: 'edusphere-dev',
    plan: 'enterprise',
    maxUsers: 1000,
    settings: {
      defaultLanguage: 'he',
      allowedLlmProviders: ['ollama', 'openai', 'anthropic'],
      maxStorageGb: 100,
    },
  }).onConflictDoNothing();

  // 2. Create test users
  await db.insert(users).values([
    {
      id: ADMIN_USER_ID,
      tenantId: DEFAULT_TENANT_ID,
      keycloakUserId: 'dev-admin-001',
      email: 'admin@edusphere.dev',
      displayName: 'Admin User',
      role: 'org_admin',
    },
    {
      id: STUDENT_USER_ID,
      tenantId: DEFAULT_TENANT_ID,
      keycloakUserId: 'dev-student-001',
      email: 'student@edusphere.dev',
      displayName: 'Test Student',
      role: 'student',
    },
  ]).onConflictDoNothing();
}
```

```

// 3. Create sample course
const [course] = await db.insert(courses).values({
  tenantId: DEFAULT_TENANT_ID,
  creatorId: ADMIN_USER_ID,
  title: 'Introduction to Semantic Learning',
  description: 'A sample course demonstrating EduSphere capabilities',
  isPublished: true,
  tags: ['demo', 'tutorial'],
}).returning();

// 4. Create sample module
const [mod] = await db.insert(modules).values({
  tenantId: DEFAULT_TENANT_ID,
  courseId: course!.id,
  title: 'Module 1: Getting Started',
  orderIndex: 0,
}).returning();

console.log('✓ Seed complete');
console.log(` Tenant: ${DEFAULT_TENANT_ID}`);
console.log(` Admin: ${ADMIN_USER_ID}`);
console.log(` Course: ${course!.id}`);
}

seed().catch(console.error);

```

15. TypeScript Graph Query Helpers

typescript

```

// File: packages/db/src/graph/client.ts

import { sql, type PostgresJsDatabase } from 'drizzle-orm';
import { GRAPH_NAME } from '../schema/_shared.ts';

/**
 * Execute a Cypher query through Apache AGE.
 * Always loads AGE extension before execution.
 *
 * @example
 * ```typescript
 * const results = await executeCypher(db, {
 *   query: `MATCH (c:Concept {tenant_id: $tenant_id})
 *           WHERE c.label CONTAINS $search
 *           RETURN c.id, c.label LIMIT 10`,
 *   columns: ['id', 'label'],
 *   tenantId: ctx.tenantId,
 * });
 * ```
 */
export async function executeCypher<T extends Record<string, unknown>>(
  db: PostgresJsDatabase,
  options: {
    query: string;
    columns: string[];
    tenantId?: string;
  },
): Promise<T[]> {
  const { query, columns } = options;

  // Build RETURN type declaration
  const returnType = columns.map((col) => `${col} agtype`).join(', ');

  const result = await db.execute(sql` 
    LOAD 'age';
    SET search_path = ag_catalog, "$user", public;
    SELECT * FROM cypher(${GRAPH_NAME}, ${cypher$$ {sql.raw(query)} $cypher$}
    AS (${sql.raw(returnType)});`);

  // Parse agtype results into plain objects
  return (result.rows as any[]).map((row) => {
    const parsed: Record<string, unknown> = {};
    for (const col of columns) {
      const val = row[col];
      // agtype values come as strings, need JSON.parse for objects/arrays
    }
  });
}

```

```
try {
  parsed[col] = typeof val === 'string' ? JSON.parse(val) : val;
} catch {
  parsed[col] = val;
}
}

return parsed as T;
});

}

/***
 * Add a vertex to the knowledge graph.
 */

export async function addVertex(
  db: PostgresJsDatabase,
  label: string,
  properties: Record<string, unknown>,
): Promise<void> {
  const propsStr = Object.entries(properties)
    .map(([k, v]) => `${k}: ${JSON.stringify(v)}`)
    .join(', ');

  await db.execute(sql`  

    LOAD 'age';  

    SET search_path = ag_catalog, "$user", public;  

    SELECT * FROM cypher(${GRAPH_NAME}, $$  

      CREATE (:${sql.raw(label)} ${sql.raw(propsStr)})  

      $$) AS (v agtype);  

`);
}

/***
 * Add an edge between two vertices.
 */

export async function addEdge(
  db: PostgresJsDatabase,
  fromLabel: string,
  fromId: string,
  toLabel: string,
  toId: string,
  edgeLabel: string,
  properties: Record<string, unknown> = {},
): Promise<void> {
  const propsStr = Object.entries(properties)
    .map(([k, v]) => `${k}: ${JSON.stringify(v)}`)
    .join(', ');
}
```

```

const edgeProps = propsStr ? ` ${propsStr}` : "";

await db.execute(sql`  

LOAD 'age';  

SET search_path = ag_catalog, "$user", public;  

SELECT * FROM cypher(${GRAPH_NAME}, $$  

MATCH (a:${sql.raw(fromLabel)} {id: '${sql.raw(fromId)}'})  

(b:${sql.raw(toLabel)} {id: '${sql.raw(toId)}'})  

CREATE (a)-[:${sql.raw(edgeLabel)}${sql.raw(edgeProps)}]->(b)  

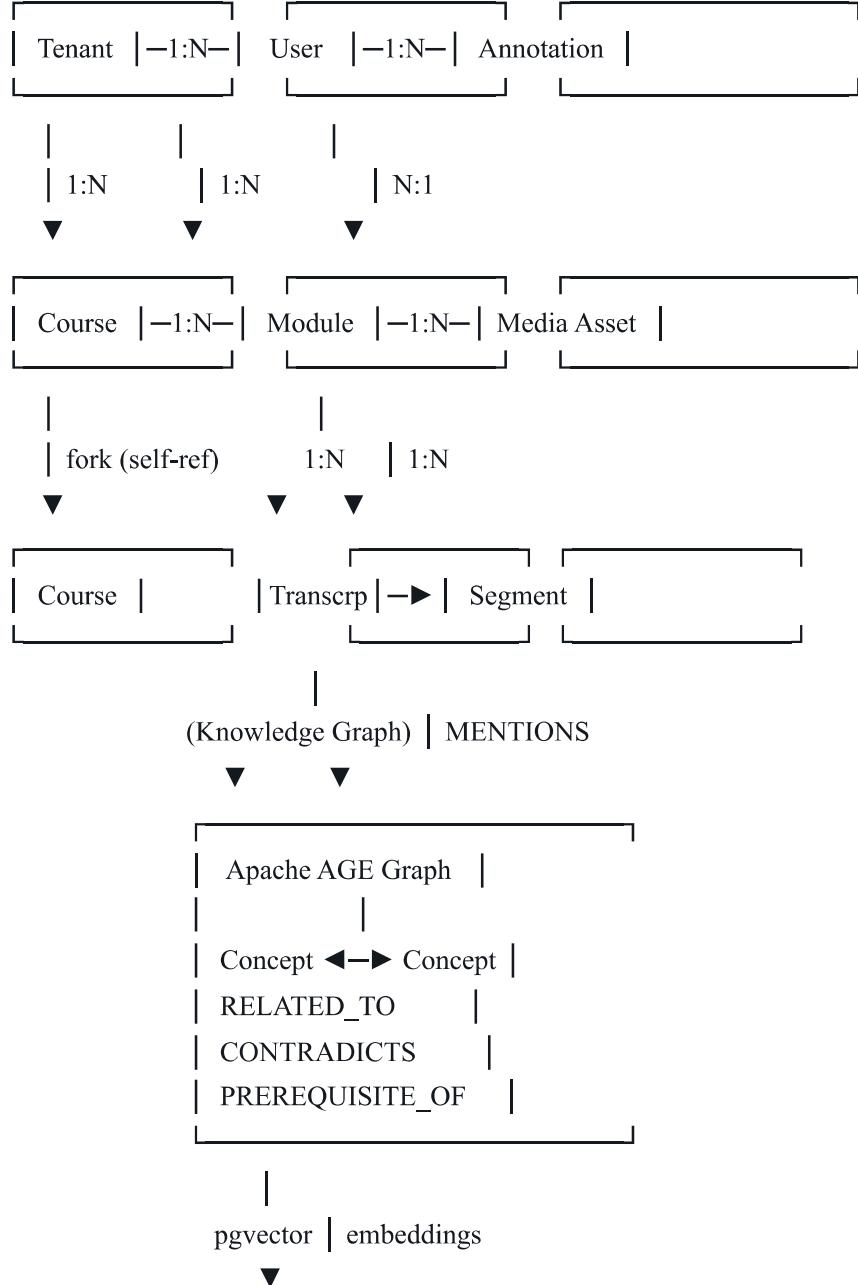
$$) AS (e agtype);  

`);

}

```

Entity Relationship Summary



Embedding Tables	
content_embeddings	
concept_embeddings	
annotation_embeds	

Total Table Count: 16

#	Table	Tenant-Isolated	RLS
1	(tenants)	N/A (IS tenant)	✓
2	(users)	✓	✓
3	(courses)	✓	✓
4	(modules)	✓	✓
5	(media_assets)	✓	✓
6	(transcripts)	✓	✓
7	(transcript_segments)	✓	✓
8	(annotations)	✓	✓
9	(collab_documents)	✓	✓
10	(crdt_updates)	✓	✓
11	(collab_sessions)	✓	✓
12	(agent_definitions)	✓	✓
13	(agent_executions)	✓	✓
14	(content_embeddings)	✓	✓
15	(annotation_embeddings)	✓	✓
16	(concept_embeddings)	✓	✓